

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/272520403>

Active inference and epistemic value

Article in *Cognitive Neuroscience* · February 2015

DOI: 10.1080/17588928.2015.1020053 · Source: PubMed

CITATIONS

588

READS

1,581

6 authors, including:



Karl J Friston

University College London

1,228 PUBLICATIONS 256,011 CITATIONS

[SEE PROFILE](#)



Dimitri Ognibene

Università degli Studi di Milano-Bicocca

89 PUBLICATIONS 1,161 CITATIONS

[SEE PROFILE](#)



Thomas FitzGerald

University of East Anglia

62 PUBLICATIONS 5,336 CITATIONS

[SEE PROFILE](#)


Some of the authors of this publication are also working on these related projects:



Hybrid Cognitive Architectures for Artificial Agents [View project](#)



Response Inhibition and Temporal Prediction Errors in Patients with Psychotic Disorders (Schizophrenia) [View project](#)



[Click for updates](#)

Cognitive Neuroscience

Publication details, including instructions for authors and subscription information:

<http://www.tandfonline.com/loi/pcns20>

Active inference and epistemic value

Karl Friston^a, Francesco Rigoli^a, Dimitri Ognibene^b, Christoph Mathys^{acd}, Thomas Fitzgerald^a & Giovanni Pezzulo^e

^a The Wellcome Trust Centre for Neuroimaging, Institute of Neurology, London, UK

^b Centre for Robotics Research, Department of Informatics, King's College London, London, UK

^c Translational Neuromodeling Unit (TNU), Institute for Biomedical Engineering, University of Zürich and ETH Zürich, Zürich, Switzerland

^d Laboratory for Social and Neural Systems Research (SNS Lab), Department of Economics, University of Zürich, Zürich, Switzerland

^e Institute of Cognitive Sciences and Technologies, National Research Council, Rome, Italy

Accepted author version posted online: 17 Feb 2015. Published online: 13 Mar 2015.

To cite this article: Karl Friston, Francesco Rigoli, Dimitri Ognibene, Christoph Mathys, Thomas Fitzgerald & Giovanni Pezzulo (2015): Active inference and epistemic value, Cognitive Neuroscience, DOI: [10.1080/17588928.2015.1020053](https://doi.org/10.1080/17588928.2015.1020053)

To link to this article: <http://dx.doi.org/10.1080/17588928.2015.1020053>

PLEASE SCROLL DOWN FOR ARTICLE

Taylor & Francis makes every effort to ensure the accuracy of all the information (the "Content") contained in the publications on our platform. However, Taylor & Francis, our agents, and our licensors make no representations or warranties whatsoever as to the accuracy, completeness, or suitability for any purpose of the Content. Any opinions and views expressed in this publication are the opinions and views of the authors, and are not the views of or endorsed by Taylor & Francis. The accuracy of the Content should not be relied upon and should be independently verified with primary sources of information. Taylor and Francis shall not be liable for any losses, actions, claims, proceedings, demands, costs, expenses, damages, and other liabilities whatsoever or howsoever caused arising directly or indirectly in connection with, in relation to or arising out of the use of the Content.

This article may be used for research, teaching, and private study purposes. Any substantial or systematic reproduction, redistribution, reselling, loan, sub-licensing, systematic supply, or distribution in any form to anyone is expressly forbidden. Terms & Conditions of access and use can be found at <http://www.tandfonline.com/page/terms-and-conditions>

Active inference and epistemic value

Karl Friston¹, Francesco Rigoli¹, Dimitri Ognibene², Christoph Mathys^{1,3,4}, Thomas Fitzgerald¹, and Giovanni Pezzulo⁵

¹The Wellcome Trust Centre for Neuroimaging, Institute of Neurology, London, UK

²Centre for Robotics Research, Department of Informatics, King's College London, London, UK

³Translational Neuromodeling Unit (TNU), Institute for Biomedical Engineering, University of Zürich and ETH Zürich, Zürich, Switzerland

⁴Laboratory for Social and Neural Systems Research (SNS Lab), Department of Economics, University of Zürich, Zürich, Switzerland

⁵Institute of Cognitive Sciences and Technologies, National Research Council, Rome, Italy

We offer a formal treatment of choice behavior based on the premise that agents minimize the expected free energy of future outcomes. Crucially, the negative free energy or *quality* of a policy can be decomposed into extrinsic and epistemic (or intrinsic) value. Minimizing expected free energy is therefore equivalent to maximizing extrinsic value or expected utility (defined in terms of prior preferences or goals), while maximizing information gain or intrinsic value (or reducing uncertainty about the causes of valuable outcomes). The resulting scheme resolves the exploration-exploitation dilemma: Epistemic value is maximized until there is no further information gain, after which exploitation is assured through maximization of extrinsic value. This is formally consistent with the Infomax principle, generalizing formulations of active vision based upon salience (Bayesian surprise) and optimal decisions based on expected utility and risk-sensitive (Kullback-Leibler) control. Furthermore, as with previous active inference formulations of discrete (Markovian) problems, ad hoc softmax parameters become the expected (Bayes-optimal) precision of beliefs about, or confidence in, policies. This article focuses on the basic theory, illustrating the ideas with simulations. A key aspect of these simulations is the similarity between precision updates and dopaminergic discharges observed in conditioning paradigms.

Keywords: Active inference; Agency; Bayesian inference; Bounded rationality; Free energy; Utility theory; Information gain; Bayesian surprise; Epistemic value; Exploration; Exploitation.

This article introduces a variational (free energy) formulation of explorative behavior and the (epistemic) value of knowing one's environment. This formulation tries to unite a number of perspectives on behavioral imperatives; namely, the exploration-exploitation dilemma and the distinction between the explicit (extrinsic) value of controlled outcomes

and their epistemic (intrinsic) value in reducing uncertainty about environmental contingencies (Bialek, Nemenman, & Tishby, 2001; Botvinick & An, 2008; Braun, Ortega, Theodorou, & Schaal, 2011; Bromberg-Martin & Hikosaka 2009; Cohen, McClure, & Yu, 2007; Daw, Niv, & Dayan, 2005; Daw, O'Doherty, Dayan, Seymour, & Dolan, 2006;

Correspondence should be addressed to: Karl Friston, The Wellcome Trust Centre for Neuroimaging, Institute of Neurology, 12 Queen Square, London, UK WC1N 3BG. E-mail: k.friston@ucl.ac.uk

The authors have no disclosures or conflict of interest.

The Wellcome Trust funded this work. KJF is funded by the Wellcome Trust [088130/Z/09/Z]. DO is funded by the European Community's Seventh Framework Programme [FP7/2007-2013] project DARWIN [Grant No: FP7-270138]. GP is funded by the European Community's Seventh Framework Programme (FP7/2007-2013) project Goal-Leaders (Grant No: FP7-ICT-270108) and the HFSP (Grant No: RGY0088/2014).

Friston et al., 2014; Pezzulo & Castelfranchi, 2009; Schmidhuber, 1991; Solway & Botvinick, 2012; Still, 2009; Tishby & Polani, 2010). In particular, it addresses how resolving uncertainty “makes the world interesting and exploitable” (Still & Precup, 2012, p. 139). It is the resolution of uncertainty that we associate with the intrinsic value of behavior, which we assume is synonymous with epistemic value. Our basic approach is to cast optimal behavior in terms of inference, where actions are selected from posterior beliefs about behavior. This allows one to frame goals and preferences in terms of prior beliefs, such that goals are subsequently fulfilled by action (Botvinick & Toussaint, 2012; Kappen, Gomez, & Opper, 2012; Toussaint & Storkey, 2006). This furnishes an explanation of behavior in terms of one straightforward imperative—to minimize surprise or, equivalently, to maximize Bayesian model evidence.

The resulting *active inference* scheme unifies conventional treatments of normative behavior under uncertainty. Classical treatments generally consider belief updates and action selection separately, calling on Bayesian inference to optimize beliefs and other schemes (such as dynamic programming) to select actions (Bonet & Geffner, 2014; Hauskrecht, 2000; Kaelbling, Littman, & Cassandra, 1998). Treating action selection as (active) inference means that both state estimation and the ensuing behavior can be described as a minimization of variational free energy or surprise.¹ In this setting, action reduces the difference between the current and (unsurprising) goal states that are defined by prior expectations, much like cybernetic formulations (Miller, Galanter, & Pribram, 1960). This difference can be reduced in two ways. First, by executing a *pragmatic action*, that fulfills goals directly (i.e., exploitation); for example, by visiting a known reward site in the context of foraging. Second, by performing an *epistemic action* (i.e., exploration) to disclose information that enables pragmatic action in the long run; for example, exploring a maze to discover unknown reward sites (Kirsh & Maglio, 1994). Clearly, most behavior has both pragmatic and epistemic aspects. Epistemic actions are the focus of much current research (Andreopoulos & Tsotsos, 2013; Ferro, Ognibene, Pezzulo, & Pirrelli, 2010; Kamar & Horvitz, 2013; Lepora, Martinez-Hernandez, & Prescott, 2013;

Lungarella & Sporns, 2006; Ognibene, Chinellato, Sarabia, & Demiris, 2013; Ognibene, Volpi, Pezzulo, & Baldassare, 2013; Pezzementi, Plaku, Reyda, & Hager, 2011; Schneider et al., 2009; Singh, Krause, Guestrin, & Kaiser, 2009). For example, the coastal navigation algorithm (Roy, Burgard, Fox, & Thrun, 1999) shows that epistemic actions can sometimes increase the distance from a goal. In this example, agents move toward a familiar location (and away from the goal) to plan a path to the goal with greater confidence. This example illustrates the exploration-exploitation dilemma encountered at every decision point—and the implicit choice between epistemic and pragmatic actions. This choice is usually addressed in the setting of reinforcement learning (Dayan, 2009; Humphries, Khamassi, & Gurney, 2012; Still & Precup, 2012) but here we place more emphasis on planning or inference (Attias, 2003; Botvinick & Toussaint, 2012). In short, we offer a solution to exploration-exploitation dilemma that rests solely on the minimization of expected free energy.

In active inference, constructs like reward, utility, epistemic value, etc. are described in terms of prior beliefs or preferences. In other words, preferred outcomes are simply outcomes one expects, a priori, to be realized through behavior (e.g., arriving at one’s destination or maintaining physiological states within some homeostatic range). This formulation of utility has a number of advantages. First, it eliminates any ad hoc parameters in classical schemes (such as softmax parameters, temporal discounting, etc.). Second, it reveals the formal relationships among classical constructs, enabling their interpretation in terms of beliefs or expectations. For example, we have previously shown that the softmax parameter in classical (utilitarian) choice models corresponds to the precision or confidence in posterior beliefs about policies—and that this precision increases with expected utility (Friston et al., 2014). Third, this formulation is equipped with a relatively simple and biologically plausible process theory based upon variational message passing (Friston et al., 2013). This can be potentially useful when looking for the neuronal correlates of message passing in decision-making paradigms. Finally, casting rewards and value as probabilistic beliefs means that intrinsic and extrinsic values share a common currency. This means one can express (extrinsic) reward in terms of (epistemic) information gain and quantify their relative contributions to behavior.

Formally speaking, we resolve the exploration-exploitation dilemma by endowing agents with prior beliefs that they will minimize the expected free

¹Variational free energy was introduced by Richard Feynman to solve inference problems in quantum mechanics and can be regarded as a generalization of thermodynamic free energy. In this paper, free energy refers to variational free energy. We will see later that minimizing free energy (or maximizing negative free energy) corresponds to maximizing expected value.

energy of future outcomes. In other words, the agent will be surprised if it behaves in a way that is not Bayes optimal. Because expected free energy determines action selection, the resulting behavior is necessarily Bayes optimal. Crucially, expected free energy is minimized over an extended timescale, making exploration a necessary and emergent aspect of optimal behavior. Expected free energy can be expressed as the Kullback-Leibler (KL) divergence between the posterior (predictive) and prior (preferred) distributions over future outcomes, plus the expected entropy of those observations, given their causes. In brief, minimizing this divergence ensures preferred outcomes are actively sampled from the environment, while minimizing the expected entropy resolves uncertainty about the (hidden) states causing those outcomes.² Intuitively, these two aspects of emergent behavior (sampling preferred outcomes and minimizing expected uncertainty) correspond to exploitation and exploration, respectively. Interestingly, the negative expected free energy can also be expressed as the expected divergence between the posterior (predictive) distribution over hidden states with and without future observations, plus the expected utility (defined as the log of the prior probability of future states). We will associate these terms with epistemic and extrinsic value respectively.

We have shown previously that minimizing the divergence between the posterior predictive distribution and prior preferences produces behavior that is risk-sensitive or KL optimal (Friston et al., 2013). Here, we show that this risk-sensitive control is a special case of minimizing expected free energy, which effectively supplements expected utility with a KL divergence that reflects epistemic value, mutual information, information gain, or Bayesian surprise, depending upon one's point of view. The KL divergence is also known as *relative entropy* or *information gain*. This means that minimizing expected free energy maximizes information gain (Ognibene & Demiris, 2013; Sornkarn, Nanayakkara, & Howard, 2014; Tishby & Polani, 2010) or, heuristically, satisfies curiosity by reducing uncertainty about the world (Schmidhuber, 1991). An alternative perspective on this epistemic quantity is afforded by *Bayesian surprise*; namely, the KL divergence between prior and posterior beliefs (Bruce & Tsotsos, 2009; Itti & Baldi, 2009).

²Note the dialectic between minimizing the entropy expected in the future and maximizing the entropy of current beliefs—implicit in minimizing free energy Friston et al. (2012). "Perceptions as hypotheses: Saccades as experiments." *Front Psychol.* 3: 151.

However, in this case, the Bayesian surprise pertains to future states that have yet to be observed.

For readers who are familiar with our previous work on active inference, this paper introduces a generic formulation that combines earlier work on optimal choice behavior (Friston et al., 2014) with formulations of salience based on sampling the world to resolve uncertainty (Friston, Adams, Perrinet, & Breakspear, 2012). These two formulations can be regarded as special cases of minimizing expected free energy, when sensory cues are unambiguous and when outcomes have only epistemic value, respectively. In this article, we show that minimizing expected free energy provides an inclusive perspective on several other established formulations of behavior.

In what follows, we introduce the basic formalism behind active inference, with a special focus on epistemic value and how this emerges under active (Bayesian) inference. The second section considers (biologically plausible) variational message passing schemes that can be used to simulate active inference in the context of partially observed Markov decision processes (Kaelbling et al., 1998) or to model empirical choice behavior. The final sections present simulations of exploration and exploitation, using a simple foraging game to illustrate the fundamental role of epistemic value in actively resolving uncertainty about goal-directed behavior. These sections consider planning and learning as inference, respectively. In future work, we will consider the optimization of models per se in terms of Bayesian model selection (structure learning) and the role of Bayesian model averaging in contextualizing shallow (model-free) and deep (model-based) models.

ACTIVE INFERENCE

This section describes active inference, in which inference and behavior are seen as consequences of minimizing variational free energy or, equivalently, maximizing Bayesian model evidence (Friston, 2010). We have previously considered epistemic value and salience using continuous time predictive coding schemes and saccadic searches (Friston et al., 2012). Here, we will use a discrete time and state space formulation of Bayes optimal behavior to show that information gain is a necessary consequence of minimizing expected free energy.

This formulation rests upon two key distinctions. First, we distinguish between a real world process that generates observations and an agent's internal model

of that process. These are referred to as the *generative process* and *generative model* respectively. The process and model are coupled in two directions: (sensory) observations generated by the generative process are observed by the agent, while the agent acts on the world to change that process. We will see that action serves to minimize the same quantity (variational free energy) used to make inferences about the hidden causes of observations. Crucially, *action* is a real variable that acts on the generative process, while the corresponding hidden cause in the generative model is a *control state*. This means the agent has to infer its behavior by forming beliefs about control states, based upon the observed consequences of its action.

We will adopt the formalism of partially observed Markov decision processes (POMDP). This is just a way of describing transitions among (discrete) states, under the assumption that the probability of the next state depends on, and only on, the current state. The partially observed aspect of the ensuing Markovian process means that the states of the generative process are hidden and have to be inferred through a limited set of (possibly noisy) observations.³

Notation: We use conventional notation, where the parameters of categorical distributions over discrete states $s \in S \in \{1, \dots, J\}$ are denoted by $J \times 1$ vectors of expectations $\hat{S} \in [0, 1]$, while the \sim notation denotes sequences of variables over time. The entropy of a probability distribution $P(a) = \Pr(A = a)$ will be denoted by $H(A) = H[P(a)] = E_{P(a)}[-\ln P(a)]$ and the relative entropy by the Kullback-Leibler (KL) divergence $D[Q(a)||P(a)] = E_{Q(a)}[\ln Q(a) - \ln P(a)]$. The dot notation means $A \cdot B = A^T B$ and $A \times B$ denotes the Hadamard (or element by element) product of two matrices. Similarly, $\ln A$ denotes the logarithm of the elements of a matrix.

Definition: Active inference rests on the tuple $(P, Q, R, S, A, U, \Omega)$:

- A finite set of observations Ω
- A finite set of actions A
- A finite set of hidden states S
- A finite set of control states U

- A *generative process* over observations $\tilde{o} \in \Omega$, hidden states $\tilde{s} \in S$, and action $\tilde{a} \in A$

$$R(\tilde{o}, \tilde{s}, \tilde{a}) = \Pr(\{o_0, \dots, o_t\} = \tilde{o}, \{s_0, \dots, s_t\} = \tilde{s}, \{a_0, \dots, a_{t-1}\} = \tilde{a})$$

- A *generative model* over observations $\tilde{o} \in \Omega$, hidden $\tilde{s} \in S$, and control $\tilde{u} \in U$ states $P(\tilde{o}, \tilde{s}, \tilde{u}|m) = \Pr(\{o_0, \dots, o_T\} = \tilde{o}, \{s_0, \dots, s_T\} = \tilde{s}, \{u_t, \dots, u_T\} = \tilde{u})$, with parameters θ
- An *approximate posterior* over hidden and control states such that $Q(\tilde{s}, \tilde{u}) = \Pr(\{s_0, \dots, s_T\} = \tilde{s}, \{u_t, \dots, u_T\} = \tilde{u})$ with parameters or expectations $(\hat{s}, \hat{\pi})$, where $\pi \in \{1, \dots, K\}$ is a policy that indexes a sequence of control states $(\tilde{u}|\pi) = (u_t, \dots, u_T|\pi)$

Remark: The *generative process* describes the environment in terms of transitions among hidden states that generate observed outcomes. These transitions depend upon actions, which are sampled from approximate posterior beliefs about control states. In turn, these beliefs are formed using a *generative model* (denoted by m) of how observations are generated. The generative model describes what the agent believes about the world, where (approximate posterior) beliefs about hidden states and control states are encoded by expectations. This is a slightly unusual setup because there is a distinction between actions (that are part of a generative process) and control states (that are part of the generative model). This distinction allows actions to be sampled from posterior beliefs about control, effectively converting an optimal control problem into an optimal inference problem (Attias, 2003; Botvinick & Toussaint, 2012). Furthermore, note that (unlike the generative process) the generative model includes beliefs about future states.

So far, we have just described the agent's world and its model of that world. To describe the agent's exchange with its environment, we have to specify how its expectations depend upon observations and how its action depends upon expectations. In other words, we have to close the perception-action cycle (Fuster, 2004). In brief, we will make one assumption; namely, that both actions and expectations minimize the free energy of observations. More precisely, we will assume that expectations minimize free energy and the expectations of control states prescribe action at the current time t :

³For readers interested in technical details, the simulations (and figures) reported in this paper can be reproduced by downloading the academic freeware SPM. Annotated Matlab scripts can then be accessed through a graphical user interface (invoked by typing DEM and selecting "epistemic value"). Please visit <http://www.fil.ion.ucl.ac.uk/spm/software/>

$$\begin{aligned}
(\hat{s}^*, \hat{\pi}^*) &= \arg \min F(\tilde{o}, \tilde{s}, \tilde{\pi}) \\
\Pr(a_t = u_t) &= Q(u_t | \hat{\pi}^*) \\
F(\tilde{o}, \tilde{s}, \tilde{\pi}) &= E_Q[-\ln P(\tilde{o}, \tilde{s}, \tilde{u} | m)] - H[Q(\tilde{s}, \tilde{u})] \\
&= -\ln P(\tilde{o} | m) + D[Q(\tilde{s}, \tilde{u}) || P(\tilde{s}, \tilde{u} | \tilde{o})]
\end{aligned} \tag{1}$$

Heuristically, at each decision point or cycle the agent first figures out which states are most likely by optimizing its expectations with respect to free energy (using the generative model). After optimizing its posterior beliefs, an action is sampled from the posterior probability distribution over control states. Given this action, the environment generates a new observation (using the generative process) and a new cycle begins.

The first expression for free energy in Equation 1 shows that it is an expected energy, under the generative model, minus the entropy of the approximate posterior. This expression can be rearranged to give the second expression, which shows that free energy is an upper bound on the negative logarithm of Bayesian model evidence $-\ln P(\tilde{o} | m)$, which is also known as *surprise* or *surprisal*. The free energy is an upper bound on surprise because the divergence term cannot be less than zero (Beal, 2003). Therefore, minimizing free energy corresponds to minimizing the divergence between the approximate and true posterior. This formalizes the notion of approximate Bayesian inference in psychology and machine learning (Dayan & Hinton, 1997; Dayan, Hinton, & Neal, 1995; Helmholtz, 1866/1962). Minimizing surprise provides a nice perspective on perception which, in this setting, corresponds to updating expectations about hidden states of the world in a Bayes optimal fashion. But what about action? If action is sampled from beliefs about control states, then the agent must believe its actions will minimize free energy. We now look at this more closely.

In active inference, agents do not just infer hidden states but actively sample outcomes that minimize free energy. The aim here is to explain how agents restrict themselves to a small number of preferred outcomes (i.e., goals). This is fairly straightforward to explain if agents minimize surprise, while a priori expecting to attain their goals. More formally, if actions depend upon posterior beliefs, then actions depend on prior beliefs. This means prior beliefs entail goals because they specify action. In turn, the generative model entails prior beliefs because it comprises the likelihood over observations, an empirical prior over state transitions and a prior over

control states. These correspond to the three marginal distributions of the generative model: $P(\tilde{o}, \tilde{s}, \tilde{u} | m) = P(\tilde{o} | \tilde{s})P(\tilde{s} | \tilde{u})P(\tilde{u} | m)$. Crucially, the only self-consistent prior beliefs an agent can entertain about control states is that they will minimize free energy.⁴ One can express this formally by associating the prior probability of a policy with the path integral (from the current to the final state) of free energy expected under that policy (c.f., Hamilton's principle of least action and Feynman's path integral formulation of quantum mechanics). We will call this the quality, value, or the expected (negative) free energy of a policy, denoted by $\mathbf{Q}(\tilde{u} | \pi) := \mathbf{Q}(\pi)$:

$$\begin{aligned}
\ln P(\tilde{u} | \gamma) &= \gamma \cdot \mathbf{Q}(\pi) = \gamma \cdot (\mathbf{Q}_{t+1}(\pi) + \dots + \mathbf{Q}_T(\pi)) \\
\mathbf{Q}_t(\pi) &= E_{Q(o_t, s_t | \pi)}[\ln P(o_t, s_t | \pi)] + H[Q(s_t | \pi)]
\end{aligned} \tag{2}$$

This expression says that a policy is a priori more likely if the policy has a high quality or its expected free energy is small. Heuristically, this means that agents believe they will pursue policies that minimize the expected free energy of outcomes and implicitly minimize their surprise about those outcomes. Equivalently, policies that do not minimize expected free energy are a priori surprising and will be avoided. Put simply, not only do agents minimize free energy or surprise (Equation 1) but they also believe they will minimize free energy or surprise (Equation 2). These beliefs (Equation 2) are realized through active inference because agents minimize surprise (Equation 1). This self-consistent recursion leads to behavior that is apparently purposeful, in the sense that it appears to avoid surprising states.

The expected free energy is the free energy of beliefs about the future (not the free energy of future beliefs). More formally, the expected free energy is the energy of counterfactual outcomes and their causes expected under their posterior predictive distribution, minus the entropy of the posterior

⁴This is a fairly subtle assertion that lies at the heart of active inference. Put simply, agents will adjust their expectations to minimize the free energy associated with any given observations. However, when the agent actively samples observations, it has the opportunity to choose observations that minimize free energy—an opportunity that is only realized when the agent believes this is how it behaves. A more formal proof by *reductio ad absurdum*—that appeals to random dynamical systems—can be found in Friston and Mathys (2015). I think therefore I am. *Cognitive Dynamic Systems*. S. Haykin, IEEE press: in press. In brief, to exist, an ergodic system must place an upper bound on the entropy of its states, where entropy is the long-term average of surprise. Therefore, any system that does not (believe it will) minimize the long-term average of surprise does not (believe it will) exist.

predictive distribution over hidden states. The posterior predictive distributions are distributions over future states at $\tau > t$ expected under current beliefs: $Q(o_\tau, s_\tau | \pi) = E_{Q(s_t)}[P(o_\tau, s_\tau | s_t, \pi)]$. Notice that this predictive posterior includes beliefs about future outcomes and hidden states, while the current posterior $Q(s_t)$ just covers hidden states. In this setup, $\gamma \in \theta$ plays the role of a sensitivity or inverse temperature parameter that corresponds to the precision of, or confidence in, prior beliefs about policies.

Note that we have introduced a circular causality by specifying prior beliefs in this way: Prior beliefs about control states depend upon (approximate posterior predictive) beliefs about hidden states, which depend on observations. This means that prior beliefs about policies depend upon past observations. Indeed, we will see later that if the precision parameter γ was known, the prior and posterior beliefs would be identical. However, when the precision is a free (hyper) parameter, posterior beliefs become the prior beliefs expected under posterior precision. This may sound rather complicated but the important role of posterior precision or confidence will become increasingly evident. In brief, by making precision a free parameter, it can be optimized with respect to free energy or model evidence (unlike the inverse temperature parameter of conventional models). We now try to unpack these beliefs about policies in terms of established formulations of goal-directed behavior.

Although Equation 2 has a relatively simple form, it is not easy to see the behaviors it produces. However, with some straightforward rearrangement, two intuitive terms reveal themselves; namely, extrinsic and epistemic value (see also [Appendix A](#)).

$$\begin{aligned} Q_\tau(\pi) &= E_{Q(o_\tau, s_\tau | \pi)}[\ln P(o_\tau, s_\tau | \pi) - \ln Q(s_\tau | \pi)] \\ &= E_{Q(o_\tau, s_\tau | \pi)}[\ln Q(s_\tau | o_\tau, \pi) + \ln P(o_\tau | m) \\ &\quad - \ln Q(s_\tau | \pi)] = \underbrace{E_{Q(o_\tau | \pi)}[\ln P(o_\tau | m)]}_{\text{Extrinsic value}} \\ &\quad + \underbrace{E_{Q(o_\tau | \pi)}[D[Q(s_\tau | o_\tau, \pi) || Q(s_\tau | \pi)]]}_{\text{Epistemic value}} \end{aligned} \quad (3)$$

Here, the generative model of future states $P(o_\tau, s_\tau | \pi) = Q(s_\tau | o_\tau, \pi)P(o_\tau | m)$ comprises the predictive posterior and prior beliefs about future outcomes. Note, the generative model of future states is not the generative model of states in the future, when the *predictive* posterior becomes the future posterior and the generative model of the future becomes the future generative model $P(o_\tau, s_\tau | \pi) = P(s_\tau | o_\tau, \pi)P(o_\tau | m)$. Equation 3 shows that under the generative model of

the future, the quality of a policy can be expressed in terms of extrinsic and epistemic value:

Extrinsic value: Extrinsic value is the utility $C(o_\tau | m) = \ln P(o_\tau | m)$ of an outcome expected under the posterior predictive distribution. It is this utility that encodes the preferred outcomes that lend behavior its goal-directed nature. In other words, agents consider outcomes with low utility surprising, irrespective of the policy. This means that agents (believe they) will maximize expected utility to ensure preferred outcomes. Note that, by definition, the utility of an outcome is not a function of the policy. This means the agent believes all (unsurprising) policies lead to the same preferred outcomes or goals. The degree to which expected utility dominates prior beliefs about policies rests on the precision of prior preferences. In the absence of precise goals, epistemic or intrinsic value will come to dominate policy selection.

Epistemic value: Epistemic value is the expected information gain under predicted outcomes. In other words, it reports the reduction in uncertainty about hidden states afforded by observations. Because the KL divergence (or information gain) cannot be less than zero, the information gain is smallest when the posterior predictive distribution is not informed by new observations. Heuristically, this means valuable policies will search out observations, cues or “signs” that resolve uncertainty about the state of the world (e.g., foraging to resolve uncertainty about the hidden location of food or fixating on informative part of a face to identify someone). However, when there is no posterior uncertainty, and the agent is confident about the state of the world, there can be no further information gain and epistemic value will be the same for all policies. In this case, extrinsic value will dominate policy selection.

Relationship to established formalisms

The Infomax principle: Epistemic or intrinsic value fits comfortably with a number of formulations from the visual sciences and information theory. As discussed (using continuous time formulations) in Friston et al. (2012), minimizing uncertainty about hidden states necessarily entails an increase in the mutual information between (sensory) outcomes and their (hidden) causes. Formally, this can be seen with a simple rearrangement of epistemic value to show that it is equivalent to the mutual information between hidden states and outcomes, under the posterior predictive distribution:

$$\begin{aligned}
& \underbrace{E_{Q(o_\tau|\pi)}[D[Q(s_\tau|o_\tau, \pi) || Q(s_\tau|\pi)]]}_{\text{Epistemic value}} \\
&= \underbrace{D[Q(s_\tau, o_\tau|\pi) || Q(s_\tau|\pi) Q(o_\tau|\pi)]}_{\text{Predictive mutual information}} \quad (4)
\end{aligned}$$

This means that policies with epistemic value render observations more informative about their causes. This is one instance of the Infomax principle (Linsker, 1990), which is closely related to the principle of maximum mutual information, or minimum redundancy (Barlow, 1961, 1974; Bialek et al., 2001; Najemnik & Geisler, 2005; Oja, 1989; Olshausen & Field, 1996; Optican & Richmond, 1987).

Bayesian surprise: Epistemic value is also the Bayesian surprise expected under counterfactual outcomes. Bayesian surprise is a measure of salience and is the KL divergence between a posterior and prior distribution (Itti & Baldi, 2009). Empirically, people tend to direct their gaze toward salient visual features with high Bayesian surprise (Itti & Baldi, 2009). In the current setup, the expected Bayesian surprise, or salience, is the epistemic value of a particular policy that samples (sensory) outcomes. Although the value of a policy includes Bayesian surprise, it also comprises expected utility, which contextualizes the influence of salience. In other words, salience will only drive epistemic sampling of salient information if the epistemic value of that sampling is greater than the extrinsic value of an alternative behavior. We will see examples of this later.

Value of information: The value information is the amount an agent would pay to obtain information pertaining to a decision (Howard, 1966; Krause & Guestrin, 2005; Kamar & Horvitz, 2013). In this formulation, information has no epistemic value per se but only relative to choices or policy selection; in other words, information that does not affect a choice has no value. The value of information is generally intractable to compute for complex (e.g., nonstationary) environments. Here, we offer a formulation that contextualizes the value of information (epistemic value) in relation to extrinsic value and provides a tractable (approximate Bayesian inference) scheme for its evaluation.

KL control: Optimal control problems can generally be expressed as minimizing the KL divergence between the preferred and predictive distribution over outcomes. The general idea behind KL control is to select control states that minimize the difference between predicted and desired outcomes, where the difference is measured in terms of the KL

divergence between the respective probability distributions. Minimizing this divergence is a cornerstone of risk-sensitive control (Van Den Broek, Wiegerinck, & Kappen, 2010) and utility-based free energy treatments of bounded rationality (Ortega & Braun, 2011, 2013). In the current context, risk-sensitive (KL) control can be seen as a special case of minimizing expected free energy, when outcomes unambiguously specify hidden states. In other words, when the generative process is completely observable, we can associate each outcome with a hidden state such that $o_\tau = s_\tau$ and:

$$\begin{aligned}
Q_\tau(\pi) &= E_{Q(s_\tau|\pi)}[\ln P(s_\tau|\pi) - \ln Q(s_\tau|\pi)] \\
&= - \underbrace{D[Q(s_\tau|\pi) || P(s_\tau|\pi)]}_{\text{KL divergence}} = \underbrace{E_{Q(s_\tau|\pi)}[\ln P(s_\tau|m)]}_{\text{Extrinsic value}} \\
&\quad + \underbrace{H[Q(s_\tau|\pi)]}_{\text{Epistemic value}} \quad (5)
\end{aligned}$$

In this special case, minimizing free energy minimizes the divergence between the posterior predictive distribution over states and the prior predictive distribution encoding goals. Here, the extrinsic value now becomes an expected utility over *states* and the epistemic value becomes the novelty or (posterior predictive) entropy over future states. The difference between maximizing the entropy (novelty) and relative entropy (information gain) distinguishes risk-sensitive (KL) control from free energy minimization. Only minimizing free energy allows epistemic value to guide explorative behavior in a way that fully accommodates uncertainty about a partially observed world. This can be seen clearly with a final rearrangement of the expression for the quality of a policy (see [Appendix A](#)):

$$\begin{aligned}
Q_\tau(\pi) &= E_{Q(o_\tau, s_\tau|\pi)}[\ln Q(o_\tau|s_\tau, \pi) \\
&\quad + \ln P(o_\tau|m) - \ln Q(o_\tau|\pi)] \\
&= - \underbrace{E_{Q(s_\tau|\pi)}[H[P(o_\tau|s_\tau)]]}_{\text{Predicted uncertainty}} - \underbrace{D[Q(o_\tau|\pi) || P(o_\tau|m)]}_{\text{Predicted divergence}} \quad (6)
\end{aligned}$$

This equality expresses the value of a policy in terms of the posterior predictive distribution over outcomes, as opposed to hidden states. In this formulation, expected free energy corresponds to the expected entropy or uncertainty over outcomes, given their causes, plus the KL divergence between the posterior predictive and preferred distributions. In

other words, minimizing expected free energy minimizes the divergence between predicted and preferred outcomes (i.e., predicted divergence) and any uncertainty afforded by observations (i.e., predicted uncertainty). Heuristically, this ensures observations are informative. For example, an agent who wants to avoid bright light will move to the shade, as opposed to closing its eyes. If outcomes are always informative, we revert to risk-sensitive (KL) control, expressed in terms of preferences over outcomes, as opposed to states.

In our previous formulations of active inference and risk-sensitive (KL) control, we only considered scenarios in which hidden states could be observed directly. In this paper, we will illustrate the difference between risk-sensitive (KL) control and expected free energy minimization in a more realistic setting, in which hidden states can only be inferred from particular observations. In this context, we will see that risk-sensitive (KL) control is not sufficient to explain purposeful or exploratory responses to salient cues that resolve uncertainty about the environment.

Dopamine and reward prediction errors: In the next section, we will see how approximate Bayesian inference, implicit in active inference, can be implemented using a relatively simple variational message passing scheme. We have previously discussed the biological plausibility of this scheme in terms of recursive neuronal message passing (Friston et al., 2013) and have associated dopamine with the posterior precision of beliefs about control states (Friston et al., 2014). We will see later that changes in the expected (inverse) precision are identical to changes in (negative) expected value. This is potentially important because it may explain why changes in dopamine firing have been associated with reward prediction error (Schultz, 1998). However, it has a deeper implication here: If expected precision changes with expected value, then the current formulation explains why dopamine has a multilateral sensitivity to novelty (Kakade & Dayan, 2002; Krebs, Schott, Schütze, & Düzel, 2009; Wittmann, Daw, Seymour, & Dolan, 2008), salience (Berridge, 2007), expected reward (Bunzeck & Düzel, 2006; D'Ardenne, McClure, Nystrom, & Cohen, 2008; Daw & Doya, 2006; Dayan, 2009; McClure, Daw, & Montague, 2003; O'Doherty et al., 2004; Pessiglione, Seymour, Flandin, Dolan, & Frith, 2006), epistemic value (Fiorillo, Tobler, & Schultz, 2003; Redgrave & Gurney, 2006; Bromberg-Martin & Hikosaka, 2009), and affordance (Cisek, 2007; Gurney, Prescott, & Redgrave, 2001; see also Nepora & Gurney, 2012). The study of Bromberg-Martin and Hikosaka (2009) is

particularly interesting in this context because it provides direct evidence linking dopamine responses and epistemic value. The emerging perspective also fits comfortably with recent attempts to reconcile dopamine's role in the exploration-exploitation trade-off with the role of the basal ganglia in action selection, "by testing the hypothesis that tonic dopamine in the striatum, the basal ganglia's input nucleus, sets the current exploration-exploitation trade-off" (Humphries et al., 2012, p. 1).

The close relationship between expected precision and value provides an interesting perspective on the transfer of dopaminergic responses to conditioned stimuli in operant conditioning paradigms (Schultz, 1998). From the perspective of active inference, conditioned stimuli have epistemic value because they resolve uncertainty about future outcomes (unconditioned stimuli). This perspective may also provide an inferential account of blocking and latent inhibition, in the sense that if epistemic uncertainty has already been resolved by one conditioned stimulus, then no further information gain is afforded by another. We will pursue these arguments with simulations of goal-directed behavior below. The important issue here is a dual role for dopaminergic responses in reporting precision in terms of extrinsic and epistemic value. However, functionally, there is only one precision (sensitivity) parameter that applies to, and reconciles, both aspects of value. This eliminates the need for ad hoc parameters to finesse the exploration-exploitation dilemma. We will illustrate these and other points using simulations in the last two sections.

Summary

Although minimizing expected free energy corresponds to maximizing extrinsic and epistemic value, this dual maximization is a particular perspective on the underlying imperative to minimize surprise. This means that both extrinsic and epistemic value work synergistically to increase the likelihood of preferred outcomes with the minimum of uncertainty. For example, extrinsic value depends on the posterior predictive distribution over outcomes, which is only informative when the agent can be confident about the current state (c.f., the coastal navigation example above). This means epistemic uncertainty must first be resolved (by increasing epistemic value) before expected utility comes into play. At the same time, an agent should not indulge in epistemic actions, if it is sufficiently confident it can pursue a successful plan.

These considerations are especially interesting in relation to exploration and exploitation.

In summary, minimizing free energy corresponds to approximate Bayesian inference and, in active inference, choosing the least surprising outcomes. If agents model their environments, they have to entertain posterior beliefs about the control of state transitions producing outcomes. This means we have to consider posterior beliefs about control states, which rest on prior beliefs about controlled outcomes. Using the self-consistent prior that control states minimize expected free energy (“I expect to avoid surprises”), we arrive at a process theory that offers a formal definition of extrinsic and epistemic value. Furthermore, it emphasizes the fact that purposeful behavior rests upon generative models that entertain future outcomes. This formulation accommodates a number of established perspectives; namely, the Infomax principle, the notion of Bayesian surprise in reporting the salience of cues, and KL control, which generalizes risk-sensitive control and expected utility theory. In the next section we will see how this theory prescribes a computational anatomy for Bayesian belief updating that has many similarities with message passing in the brain.

GENERATIVE MODELS AND VARIATIONAL MESSAGE PASSING

The generative model

The generative model used to model the (finite horizon Markovian) processes considered below can be expressed in terms of the following likelihood and prior distributions over observations and states up to time $t \in (0, \dots, T)$ (omitting normalization constants):

$$\begin{aligned} P(\tilde{o}, \tilde{s}, \tilde{u}, \gamma | \tilde{a}, m) &= P(\tilde{o} | \tilde{s}) P(\tilde{s} | \tilde{a}) P(\tilde{u} | \gamma) P(\gamma | m) \\ P(\tilde{o} | \tilde{s}) &= P(o_0 | s_0) P(o_1 | s_1) \dots P(o_t | s_t) \\ P(\tilde{s} | \tilde{a}) &= P(s_t | s_{t-1}, a_t) \dots P(s_1 | s_0, a_1) P(s_0 | m) \\ P(\tilde{u} | \gamma) &= \sigma(\gamma \cdot \mathbf{Q}) \end{aligned} \quad (7)$$

Here, $\sigma(\cdot)$ is a softmax function. The first equality expresses the generative model in terms of the likelihood of observations given the hidden states (first term) and subsequent *empirical* prior beliefs. Empirical priors are probability distributions over unknown variables that depend on other unknown variables. Empirical priors are a universal aspect of hierarchical Bayesian models; for example, parametric empirical Bayes (Kass & Steffey, 1989).

In effect, empirical priors are informed by observations under hierarchical constraints. The likelihood in the second equality implies that observations depend on, and only on, the current hidden state. The third equality expresses (empirical) prior beliefs about state transitions. For simplicity, we assume that agents know their past actions. The final equality expresses beliefs about policies in terms of their quality or value. In short, this model represents past hidden states and future choices, under the belief that controlled transitions from the current state will minimize the expected free energy of future states.

This model can be parameterized in a fairly straightforward way, using the notation $P(o_t = i | s_t = j, \mathbf{A}) = \mathbf{A}_{ij} \Leftarrow P(o_t | s_t) = \mathbf{A}$

$$\begin{aligned} P(o_t | s_t) &= \mathbf{A} \\ P(s_{t+1} | s_t, u_t) &= \mathbf{B}(u_t) \\ P(o_\tau | m) &= \mathbf{C}_\tau \\ P(s_0 | m) &= \mathbf{D} \\ P(\gamma | m) &= \Gamma(\alpha, \beta) \end{aligned} \quad (8)$$

These equalities mean that the categorical distributions over observations, given the hidden states, are encoded by the matrix $\mathbf{A} \in \theta$ that maps from hidden states to outcomes. Similarly, the transition matrices $\mathbf{B}(u_t) \in \theta$ encode transition probabilities from one state to the next, under the current control state. The vectors $\mathbf{C} \in \theta$ and $\mathbf{D} \in \theta$ encode prior distributions over future outcomes and initial states, respectively. The priors over future outcomes specify their utility $C(o_\tau | m) = \ln P(o_\tau | m) = \ln \mathbf{C}_\tau$. Finally, the prior over precision has a standard gamma distribution, with shape and rate parameters (in this paper) $\alpha = 64$ and $\beta = 4$.

The vector \mathbf{Q} contains the values of each policy at the current time. These values can be expressed in terms of the parameters above using the expression for expected free energy in Equation (6) and Appendix A:

$$\begin{aligned} \mathbf{Q}(\pi) &= \mathbf{Q}_{t+1}(\pi) + \dots + \mathbf{Q}_T(\pi) \\ \mathbf{Q}_\tau(\pi) &= \underbrace{\mathbf{1} \cdot (\mathbf{A} \times \ln \mathbf{A}) \hat{s}_\tau(\pi)}_{\text{Predicted uncertainty}} \\ &\quad - \underbrace{(\ln \hat{o}_\tau(\pi) - \ln \mathbf{C}_\tau) \cdot \hat{o}_\tau(\pi)}_{\text{Predicted divergence}} \\ \hat{s}_\tau(\pi) &= \mathbf{B}(u_t | \pi) \dots \mathbf{B}(u_1 | \pi) \hat{s}_1 \\ \hat{o}_\tau(\pi) &= \mathbf{A} \hat{s}_\tau(\pi) \end{aligned} \quad (9)$$

Where $\hat{s}_\tau(\pi)$ are the expected states at time τ under policy π and $\mathbf{1}$ is a column vector of ones. Note that when there is no uncertainty about future states, we have $\mathbf{1} \cdot (\mathbf{A} \times \ln \mathbf{A}) \hat{s}_\tau(\pi) = \ln(\mathbf{A} \hat{s}_\tau(\pi)) \cdot \mathbf{A} \hat{s}_\tau(\pi)$ and the value of a policy depends only on expected utility $\mathbf{Q}_\tau(\pi) = \hat{o}_\tau(\pi) \cdot \ln \mathbf{C}_\tau$. In other words, policies have no epistemic value when they lead to no further information gain.

Approximate Bayesian inference

Having specified the generative model, variational Bayes now offers a generic scheme for approximate Bayesian inference that finesses the combinatoric and analytic intractability of exact inference (Beal, 2003; Fox & Roberts, 2011). Variational Bayes rests on a factorization of approximate posterior beliefs that greatly reduces the number of expectations required to encode it. The factorization we focus on exploits the Markovian nature of the generative model and has the following form (see Friston et al., 2013 for details):

$$\begin{aligned} Q(\tilde{s}, \tilde{u}, \gamma | \mu) &= Q(s_0 | \tilde{s}_0) \dots Q(s_T | \tilde{s}_T) \\ &Q(u_t, \dots, u_T | \tilde{\pi}) Q(\gamma | \tilde{\gamma}) \\ Q(\gamma | \tilde{\gamma}) &= \Gamma(\alpha, \tilde{\beta} = \alpha / \tilde{\gamma}) \end{aligned} \quad (10)$$

This assumes a factorization over hidden states, (future) control states, and precision. It is this factorization that renders the inference approximate and resolves many of the intractable problems of exact inference. For example, the factorization does not consider sequences of hidden states, which means we only have to evaluate sequences of control states (as opposed to all possible sequences of controlled state transitions). We have assumed here that the posterior marginal over precision is, like its conjugate prior, a gamma distribution. The rate parameter of this posterior belief $\tilde{\beta} = \alpha / \tilde{\gamma}$ corresponds to temperature in classic formulations. However, it is no longer a fixed parameter but a sufficient statistic of beliefs about policies.

Given the generative model (Equation 7) and the mean field assumption (Equation 10), it is straightforward to solve for the expectations that minimize variational free energy (see Appendix B).

$$\begin{aligned} \hat{s}_t &= \sigma(\ln \mathbf{A} \cdot o_t + \ln(\mathbf{B}(a_{t-1}) \hat{s}_{t-1})) \\ \tilde{\pi} &= \sigma(\tilde{\gamma} \cdot \mathbf{Q}) \\ \tilde{\gamma} &= \frac{\alpha}{\beta - \mathbf{Q} \cdot \tilde{\pi}} \end{aligned} \quad (11)$$

Iterating these self-consistent equations until convergence produces the posterior expectations that minimize free energy and provides Bayesian estimates of the unknown variables. This means that expectations change over two timescales: A fast timescale that updates posterior beliefs between observations and a slow timescale that updates posterior beliefs as new observations are sampled. We have speculated (Friston, Samothrakakis, & Montague, 2012) that these updates may be related to nested electrophysiological oscillations, such as phase coupling between gamma and theta oscillations in prefrontal–hippocampal interactions (Canolty et al., 2006). See also (Penny, Zeidman, & Burgess, 2013). The forms of these updates are remarkably simple and we now consider each in turn.

The first equation updates expectations about hidden states and corresponds to *perceptual inference* or *state estimation*. This is essentially a Bayesian filter that combines predictions based upon expectations about the previous state with the likelihood of the current observation. For simplicity, we have ignored the dependency of value on expected states that would introduce a third (optimism bias) term (see Appendix B).

The second update is just a softmax function of the value of each policy, where the sensitivity parameter or expected precision is an increasing function of expected value. This last point is quite important: It means that the sensitivity or inverse temperature, that determines the precision with which a policy is selected, increases with the expected value of those policies.

The third update optimizes expected precision. If we express these updates in terms of the posterior rate parameter, we see that changes in (inverse) precision are changes in (negative) expected value:

$\tilde{\beta} = \beta - \mathbf{Q} \cdot \tilde{\pi}$. In other words, if an observation increases the expected value of the policies entertained by an agent, then expected precision increases (i.e., temperature decreases) and the agent is implicitly more confident in selecting the next action. As noted above, this may explain why dopamine discharges have been interpreted in terms of changes in expected value (e.g., reward prediction errors). The role of the neuromodulator dopamine in encoding precision is further substantiated by noting that precision enters the

variational updates in a multiplicative or modulatory fashion. We will pursue this in the next section.

Summary

In summary, by assuming a generic (Markovian) form for the generative model, it is fairly easy to derive Bayesian updates that clarify the interrelationships between expected value and precision—and how these quantities shape beliefs about hidden states of the world and subsequent behavior. Furthermore, the anatomy of this message passing is not inconsistent with functional anatomy in the brain (see Friston et al., 2014, and Figure 1 in this paper). The implicit computational anatomy rests on reciprocal message passing between expected policies (e.g., in the striatum) and expected precision (e.g., in the substantia nigra). Expectations about policies depend upon value that, in turn, depends upon expected states of the world that are iterated forward in time—to evaluate free energy in the future (e.g., in the prefrontal cortex; Mushiake, Saito, Sakamoto, Itoyama, & Tanji, 2006) and possibly hippocampus (Pezzulo, Van der Meer, Lansink, & Pennartz, 2014). In the next section, we illustrate the basic behavior of this scheme using simulations.

INFERENCE AND PLANNING

This section considers inference using simulations of foraging for information in a relatively simple environment. Its focus is on the comparative performance when minimizing expected free energy, relative to the special cases of risk-sensitive control and maximizing expected utility or reward. In particular, we will look at the neuronal correlates of the scheme in terms of simulated dopaminergic responses. The problem we consider can be construed as searching for rewards in a T-maze. This T-maze offers primary rewards (or, in Pavlovian terms, *unconditioned stimuli*; US) such as food and cues (or *conditioned stimuli*; CS) that are not rewarding per se but disclose rewards that can be secured subsequently. The basic principles of this problem can be applied to any number of scenarios (e.g., saccadic eye movements to visual targets). This example was chosen to be as simple as possible, while illustrating a number of key points that follow from the theoretical considerations above. Furthermore, this example can also be interpreted in terms of responses elicited in reinforcement learning paradigms by unconditioned (US) and conditioned (CS) stimuli. We will call on this interpretation when relating precision updates to dopaminergic discharges.




Figure 1. This figure illustrates the cognitive and functional anatomy implied by the variational scheme in the main text. Here, we have associated the variational updates of expected hidden states with perception, of control states (policies) with action selection and, finally, expected precision with attention or salience. In terms of neuronal implementation, the requisite exchange of expectations can be likened to the exchange of neuronal signals via extrinsic connections among functionally specialized brain systems. In this (purely iconic) schematic, we have associated perception (inference about the current state of the world) with the prefrontal cortex (which plausibly interacts with the hippocampus in this context), while assigning action selection to the basal ganglia. Precision has been associated with dopaminergic projections from ventral tegmental area and substantia nigra. See main text for a full description of the equations.

The setup

The agent (e.g., rat) starts in the center of a T-maze, where either the right or left arms are baited with a reward (US). The lower arm contains a cue (CS), which tells the animal whether the reward is in the upper right or left arm. Crucially, the agent can only make two moves from any location to another (for simplicity, we do not require the agent to visit intermediate locations). Furthermore, the agent cannot leave the baited arms after they are entered. This means that the optimal behavior is to first go to the lower arm to find where the reward is located and then secure the reward at the cued location in the appropriate upper arm (i.e., the agent has to move away from the goal so that it can be secured later, as in the coastal navigation example). It is this epistemic behavior we hoped would emerge as a natural consequence of minimizing expected free energy. This may seem a remarkably simple problem but it has all the ingredients necessary to illustrate the basic aspects of behavior under active inference.

Formally, in terms of a Markov decision process, there are four control states that correspond to visiting, or sampling, the four locations (the center and three arms). For simplicity, we assume that each control state takes the agent to the associated location (as opposed to moving in a particular direction from the current location). This is analogous to place-based navigation strategies thought to be subserved by the hippocampus (e.g., Moser, Kropff, & Moser, 2008). There are four (locations) times two (right and left reward) hidden states and 16 outcomes. The 16 outcomes correspond to the four locations times four stimuli (cue right, cue left, reward, and no reward). Having specified the state space, it is now only necessary to specify the $(\mathbf{A}, \mathbf{B}, \mathbf{C}, \mathbf{D})$ matrices encoding transition probabilities and preferences. These are shown in Figure 2, where the \mathbf{A} matrix maps from hidden states to outcomes, delivering an uninformative cue at the center (first) location⁵ and a definitive cue at the lower (fourth) location. The remaining locations provide a reward (or not) with probability $a = 90\%$ depending upon the hidden context (right versus left reward).

The $\mathbf{B}(u)$ matrices encode control-dependent transitions to the corresponding location, with the exception of the baited (second and third) locations, which are hidden states that the agent cannot leave. The vector \mathbf{C} determines prior preferences about outcomes. These are expressed in terms of a softmax function of

utility, which determines the relative log probability of each outcome. Here, the utility of the rewarding stimulus is c and its absence $-c$. This means, the agent expects a rewarding outcome $\exp(2c)$ times more than a null outcome. For example, if $c = 1$ it would expect a reward about $\exp(2) \approx 8$ times more than no reward. Note that utility is always relative and has a quantitative meaning in terms of relative (log) probabilities of preferred outcomes. This is important because it endows utility with the same measure as information; namely, bits or nats (i.e., units of information or entropy, the former assuming base 2 logarithms and the latter based on natural logarithms). This highlights the close connection between value and information (see below). Finally, the vector \mathbf{D} specifies the agent's beliefs about the initial conditions; namely, that it starts at the center location with equiprobable baiting of the right or left arm.

Having specified the state space and contingencies, one can iterate the variational updates (in Equation 11 and Figure 1) to simulate behavior. In these simulations the outcomes were generated using the contingencies of the generative model. In other words, we assume the agent has already learned or optimized its model of the generative process (in terms of the model structure and its parameters). We will revisit this assumption in the last section.

Figure 3 shows the results of simulations in terms of performance (upper panel) and the dynamics of Bayesian updating in terms of precision or simulated dopaminergic responses (lower panels). The upper panel shows performance as the percentage of successful (rewarded) trials with increasing levels of utility, using six equally spaced levels from $c = 0$ to $c = 2$. Performance was assessed using 128 trials, under three different schemes: Minimizing expected free energy, risk-sensitive (KL) control, and maximizing expected utility. For completeness, we also provide the results for the free energy minimization when suppressing precision updates. The three schemes can be considered as special cases that result when successively removing terms from the expected free energy (to give reduced forms indicated by the brackets).

$$\mathbf{Q}_\tau(\pi) = E_{Q(o_\tau, s_\tau | \pi)} [\underbrace{\ln P(o_\tau | s_\tau)}_{\text{KL control}} - \underbrace{\ln Q(o_\tau | \tilde{u})}_{\text{Expected utility}} + \underbrace{\ln P(o_\tau | m)}_{\text{Expected Free energy}}] \quad (12)$$

This expression shows that risk-sensitive control is the same as minimizing expected free energy when ignoring the (predictive) entropy of outcomes given

⁵The values of one half in the first block of the \mathbf{A} matrix (Figure 2) mean that the agent cannot predict the cue from that location. In other words, there is no precise sensory information and the agent is “in the dark.”

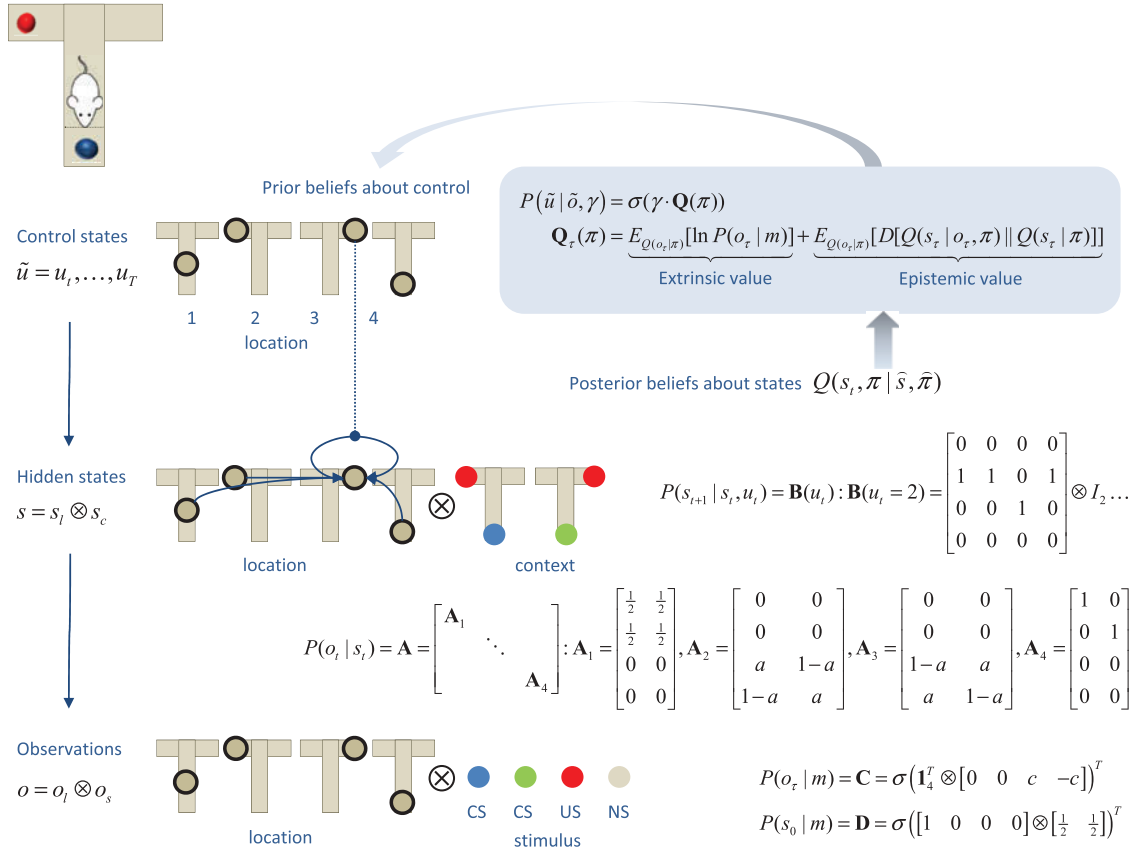


Figure 2. A schematic of the hierarchical generative model used to simulate foraging in a three-arm maze (insert on the upper left). This model contains four control states that encode movement to one of four locations (three peripheral locations and a central location). These control the transition probabilities among hidden states that have a factorial or tensor product form with two factors. The first is the location (one of four locations), while the second is one of two hidden states of the world, corresponding to a combination of cues (blue or green circles) and rewarding (red) outcomes. Each of the ensuing eight hidden states generates an observation. Some selected transitions are shown as arrows, indicating that control states attract the agent to different locations, where outcomes are sampled. The equations define the generative model in terms of its parameters $(\mathbf{A}, \mathbf{B}, \mathbf{C}, \mathbf{D}) \subset \theta$ as described in the main text. In this figure, $\sigma(\cdot)$ is a softmax function and \otimes denotes a Kronecker tensor product. Although the graphics are arranged in rows, the vectors of states are actually row vectors.

hidden states. In other words, if every hidden state generates a unique outcome, KL control and expected free energy minimization would be the same. Similarly, maximizing expected utility is the same as minimizing KL divergence, if every outcome is generated by a unique hidden state and we do not have to maximize the entropy of outcomes. In short, expected utility and classical reinforcement schemes (Sutton & Barto, 1998) are special cases of risk-sensitive control that are optimal when (and only when) different hidden states generate different outcomes. Similarly, risk-sensitive control is a special case of free energy minimization that is optimal when (and only when) different outcomes are generated by different hidden states. These special cases are important because they highlight the epistemic value of informative observations, of the sort that are precluded by noisy or context-

sensitive observations. This nesting within free energy minimization may also explain the prevalence of classical schemes in the literature, given that they generally assume hidden states are known to the agent.

The performance of the different schemes (see Figure 4, upper panel) speaks to several intuitive and useful points. First, all the schemes show an increased success rate as utility or prior preference increases; however, only expected free energy minimization attains near optimal performance (90%). One might ask why risk-sensitive control performs so poorly, given it is also sensitive to uncertainty. However, KL schemes only consider uncertainty or risk induced by many hidden states causing a single outcome, as opposed to many outcomes caused by a single state. If we had used more locations (say, with a radial maze), the benefits

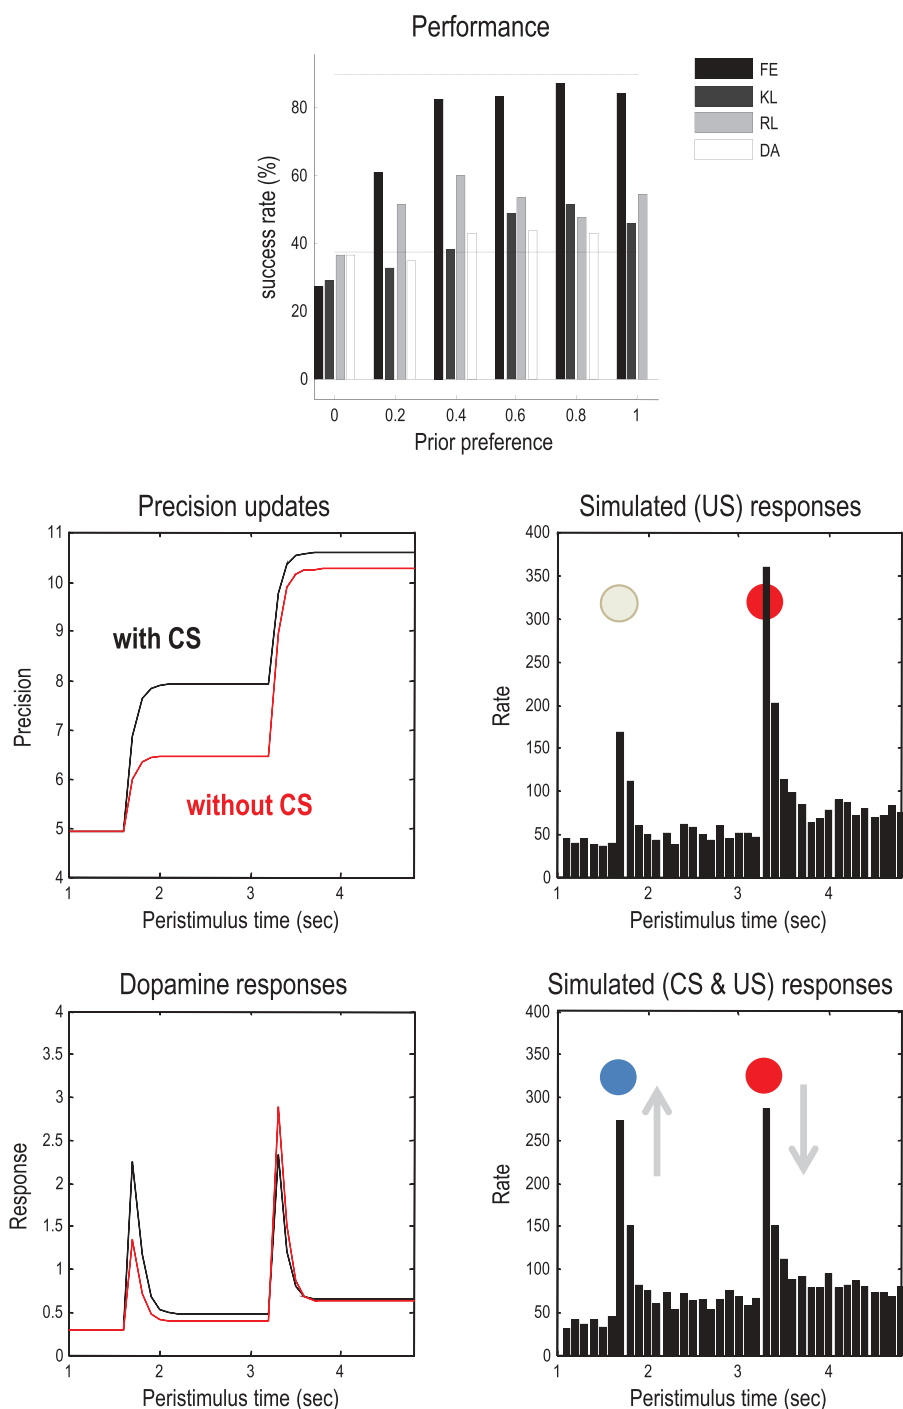


Figure 3. Upper panel: The results of 128 simulated trials assessed in terms of the probability of obtaining a reward. This performance is shown as a function of prior preference over six equally spaced levels. The four profiles correspond to active inference (FE), risk-sensitive control (KL), expected utility (RL), and active inference under fixed levels of precision (DA). See main text for a description of these schemes and how they relate to each other. The two horizontal lines show chance (bottom line) and optimal (top line) performance, respectively. Lower left panels: These report expected precision as a function of time within a trial (comprising three movements). The black lines correspond to a trial in which the cue (CS) was first accessed in the lower arm of the maze in the previous figure, after which the reward (US) was secured. The equivalent results, when staying at the center location and accessing the reward directly, are shown as red lines. The upper panel shows the expected precision and the lower panel shows simulated dopamine responses (that produce an increase in precision, which subsequently decays). Lower right panels: These show the equivalent results in terms of simulated dopamine discharges. The key thing to note here is that the responses to the cue (CS) are increased when it is informative (i.e., accessed in the lower arm), while subsequent responses to the reward (US) are decreased. See main text for details of these simulated responses.

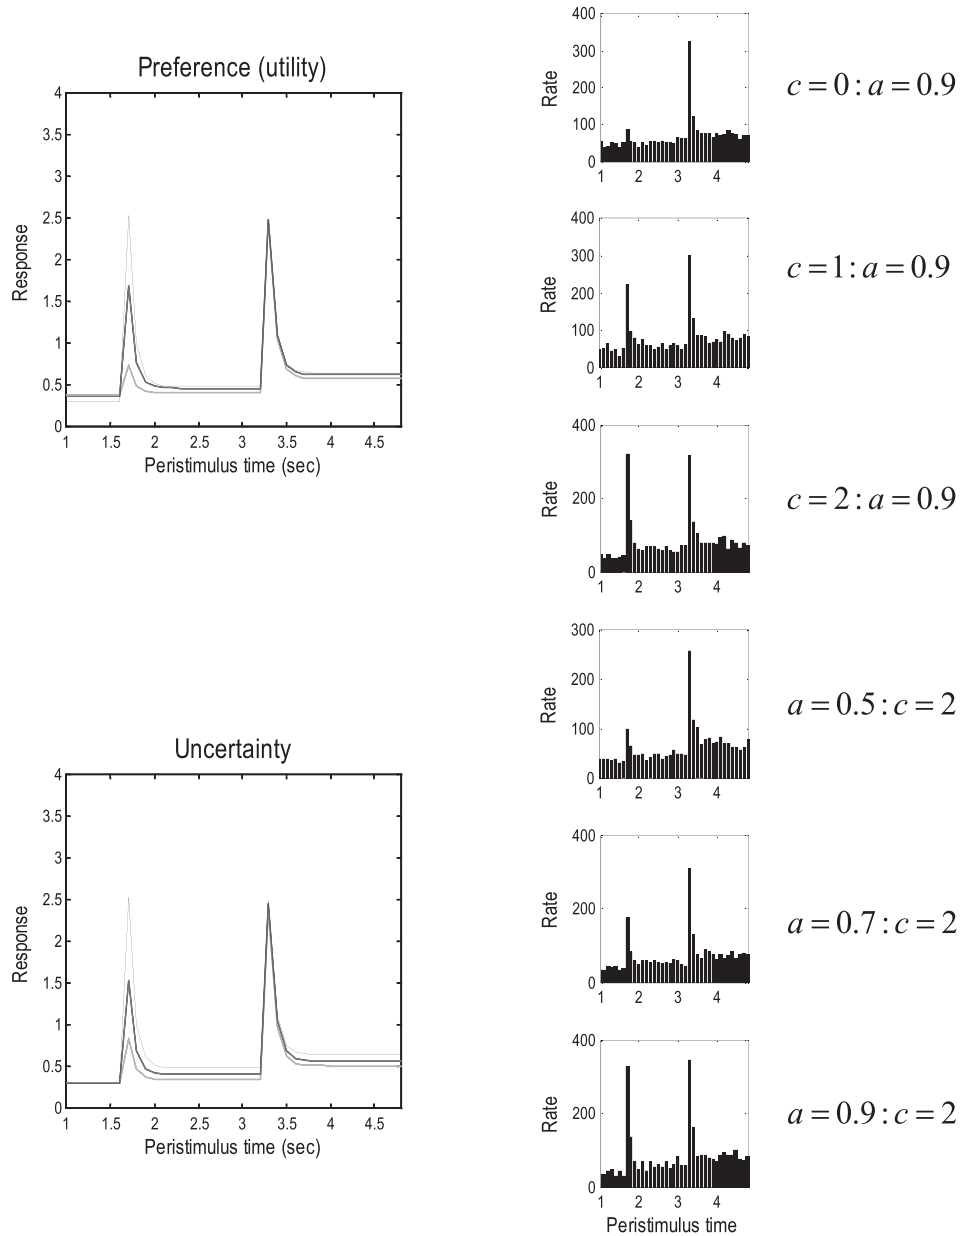


Figure 4. Simulated dopamine responses as a function of preference (upper panels) and uncertainty (lower panels). The left panels show the expected dopaminergic responses using the same format as Figure 3, for three levels of preference (utility) and uncertainty in the upper and lower panels, respectively. The right-hand panels show simulated dopaminergic firing in response to the cue (CS) and reward (US) based upon these expectations. Note that the response to the cue (CS) increases with preference and a reduction in uncertainty.

of risk-sensitive control would have been more apparent (results not shown). This follows because more locations induce more hidden states and a greater degree of uncertainty that would call for risk-sensitive control. The current setup illustrates the importance of considering both sorts of ambiguity in the mapping between causes and consequences (one-to-many and many-to-one) that calls for a minimization of expected free energy. In

this example, the most informative location is the lower arm. Visiting this location and sampling the informative cue reduces uncertainty about hidden states and enables expected utility to dominate in the second (and generally successful) move.

This optimal behavior is only apparent when utility is greater than about $c > 0.6$ nats. This brings us to our second point: If expected utility or preferences are to supervene over epistemic value, then they have to

have greater value than the information gain associated with informative outcomes. In this example, the cue resolves uncertainty about the hidden context (which arm is rewarding), thereby providing one bit or $\ln(2) = 0.6931$ nats of information. Intuitively, this means the utility must be greater than the information gain to persuade an agent to leave locations that provide unambiguous (informative) outcomes.

The third point of note is that expected utility (and risk-sensitive control) perform below chance levels when utility is zero (note that chance performance is $\frac{3}{8}$ because the agent can make two moves and is trapped by two locations). This reflects the fact that the upper arms no longer hold any utilitarian attraction and become unattractive because of the slightly ambiguous outcomes implicit in the probabilistic reward schedule. In other words, the most valuable policy is to stay in the epistemically valuable (lower) location for as long as possible.

The last set of simulations under a fixed level of precision $\hat{\gamma} = 1$ show that optimal choice behavior rests on updating expected precision, which we now look at more carefully. The lower panels of Figure 3 show how expected precision is updated during a trial of two movements and three outcomes under a high level of utility $c = 2$. These simulations are presented to highlight the similarity between precision updating and empirical dopamine responses during the presentation of conditioned and unconditioned stimuli. The upper left panel shows the expected precision over variational updates, with 16 updates between observations. The black lines correspond to a trial in which the agent accessed the conditioned stimulus (CS or cue) in the lower arm and then secured the unconditional stimulus (US or reward) on moving to an upper arm. The red lines show the equivalent updates in a second trial, when the agent stayed at the central location for the first move and was then presented with the US. In both situations, the precision increases with each successive outcome; however, the expected precision is higher in the first trial, when the CS reduces uncertainty about the hidden states or context in which the agent is operating. This reflects the greater epistemic value of accessing the cue. Crucially, the precision of the final state is roughly the same for both trials. The implication of this is that expected precision increases on presentation of the CS and, necessarily, increases less on presentation of the subsequent US, relative to presentation of the US alone.

This difference can be highlighted by plotting the expected precision in terms of simulated dopaminergic discharges, which are thought to

reflect changes in expected precision or value. More exactly, the lower left panel shows simulated dopamine discharges that would, when convolved with a decaying exponential response function (with a time constant of 16 iterations) reproduce the expected precision in the upper panel. In other words, we are assuming that dopamine mediates increases in precision that subsequently decay with a fixed time constant. In this format, one can clearly see the phasic responses of expected precision (simulated dopaminergic discharges) where, crucially, the presentation of the CS reduces the response to the US. This reproduces the well-known transfer of dopamine responses from a US to a CS in operant paradigms (Schultz, Apicella, & Ljungberg, 1993).

The right panels of Figure 3 shows simulated dopamine discharges assuming that an expected precision of one is encoded by 128 spikes per bin (and firing rates are sampled from a Poisson distribution). These are remarkably similar to empirical results, often interpreted in terms of reward prediction error and temporal difference models of value learning. However, the current framework offers a nuanced perspective; namely, the CS has epistemic value that reduces uncertainty about what will happen next. This uncertainty is already resolved when the US is presented, thereby attenuating the precision-dependent responses it elicits. Put simply, the transfer of dopaminergic responses to conditioned stimuli, in higher-order operant paradigms, can be thought of as reporting the confidence (precision) that policies will bring about predicted outcomes.

The composition of extrinsic and epistemic value implicit in expected free energy can also be used to reproduce the empirical responses of dopaminergic cells to CS under different levels of reward and uncertainty (Fiorillo et al., 2003). Figure 4 shows simulated dopamine responses under increasing utility $c = \{0, 1, 2\} : a = 0.5$ and different levels of uncertainty about the reward probability $a = \{0.5, 0.7, 0.9\} : c = 2$. In both cases, the response to the CS increases in a way that is remarkably reminiscent of empirical results (Fiorillo et al., 2003). Interestingly, the tonic responses appear to be more sensitive to uncertainty (lower panels) than utility (upper panels). This is also seen empirically, although the tonic responses reported in Fiorillo et al. (2003) increased in a ramp-like fashion under higher levels of uncertainty (i.e., $a = 0.5$). This phenomenon is not reproduced in Figure 5, however. Generally, precision increases as the trial progresses because agents become increasingly confident about their policies. One can see this general trend in Figure 4.

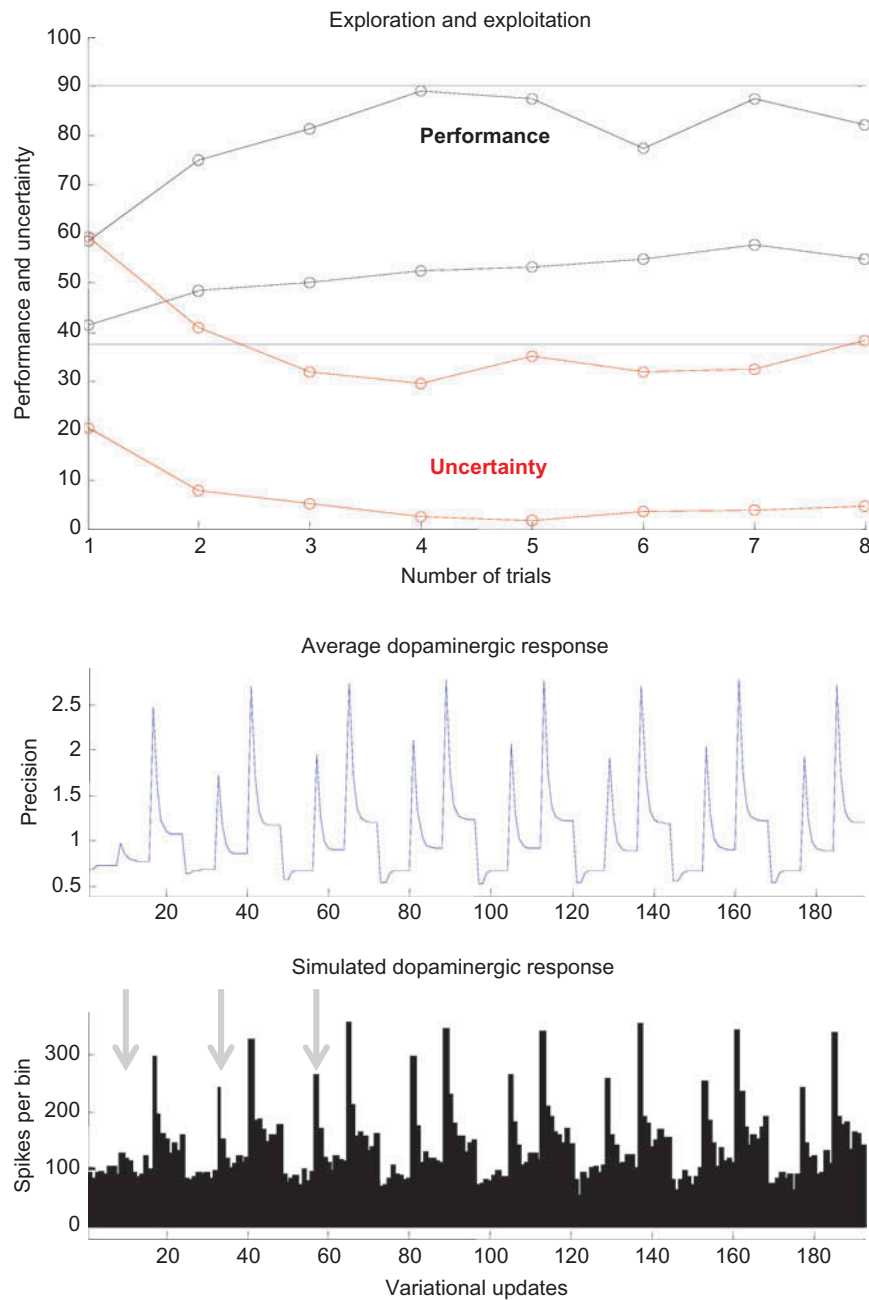


Figure 5. Upper panel: Learning in terms of success rate as a function of trials, for eight successive trials in an initially unknown maze. The results are averaged over 128 realizations. Performance (gray lines) shows a progressive improvement as uncertainty about the hidden states falls (pink lines). The equivalent performance for a conventional expected utility scheme is shown with broken lines. Lower panels: Simulated dopamine responses over all iterations and trials shown in terms of average precision (middle panel) and simulated dopaminergic spike rates (lower panel). These results demonstrate the transfer of simulated dopamine responses to the cue (CS) with learning (gray arrows).

Summary

In summary, we have seen that optimal choice behavior, in a very simple paradigm, rests on resolving uncertainty about future choices implicit in minimizing expected free energy. This aspect of

optimal behavior is clearly disclosed when decisions under uncertainty are confounded, not only by a many-to-one mapping between hidden states and outcomes, but also between outcomes and hidden states (c.f., Littman, Sutton, & Singh, 2002). In this general setting, the role of epistemic value becomes

paramount in resolving uncertainty about what to do next—a resolution that can be construed in terms of exploration or foraging for information. Furthermore, the integrative framework provided by free energy minimization enforces a dialogue between utility and information by casting both as log probabilities. This means every utility or reward can be quantified in terms of information and every bit of information has utility. We have considered the encoding of this information in terms of precision, showing that biologically plausible variational updates of expected precision are remarkably consistent with empirical dopaminergic responses. A key aspect of this formulation is that the precision of beliefs about the value of policies is itself an increasing function of expected value. This means that if dopamine reports (changes in) precision, it also reports (changes in) expected value and, implicitly, reward prediction error. So far, we have limited our discussion to planning as inference and memory. In the final section, we turn to the role of epistemic value in learning and memory, touching on some important issues that attend hierarchical inference and contextualizing behavior.

LEARNING AND MEMORY AS INFERENCE

This section uses the same setup but considers multiple trials during which the agent has to learn which locations deliver rewards and cues. In other words, we introduce an extra hierarchical level to the problem, where the hidden context now includes the mapping between locations and actions (i.e., moving to the lower arm could take it to a rewarding location). This means the agent has to learn which locations offer cues and which offer rewards. The motivation here is to illustrate a form of learning that rests on exploring the environment and to show that there is a Bayes-optimal transition from exploration to exploitation. Crucially, this solution rests upon exactly the same scheme as above—the only thing that changes is the generative model.

There are many ways of modeling learning in this context. These range from Bayesian model selection and averaging, aka structure learning (FitzGerald, Dolan, & Friston, 2014), through optimization of the model parameters ($\mathbf{A}, \mathbf{B}, \mathbf{C}, \mathbf{D}$) $\subset \theta$ with respect to expected free energy, to casting the problem as a hierarchical inference problem (c.f., Ballard, Kit, Rothkopf, & Sullivan, 2013). We will choose the latter because it requires no

extra theory⁶ and illustrates how hierarchical inference contextualizes lower-level (habitual) action selection. In brief, we will use Bayesian belief updating that embodies the prior that the mapping between locations and control states does not change from trial to trial, but the location of the reward changes between trials. The agent therefore has to learn (infer) time-invariant (contextual) aspects of its environment through exploration, before it can engage in pragmatic goal-directed behavior. We will see that this learning is an emergent property of minimizing expected free energy at each move.

The setup

Our aim was to illustrate learning as inference by introducing hierarchical uncertainty into the setup. In other words, we wanted to see if the agent could learn about its environment by introducing uncertainty about which locations offered rewards and cues. In discrete state-space formulations, hierarchical extensions involve creating product spaces, such that each lower-level state is reproduced under each level of a higher-level state. Here, we considered four higher-level hidden contexts $S^{(2)}$ corresponding to four mappings between each of the three arms of the T-maze. More specifically, we introduced four mappings between the three control states and the associated hidden location states that determine outcomes. This just involved changing the following matrices, where we denote the hierarchical level of parameters and states with superscripts (such that $\mathbf{A}, \mathbf{B}, \dots$ above become $\mathbf{A}^{(1)}, \mathbf{B}^{(1)}, \dots$):

$$\begin{aligned} \mathbf{A} &= [\mathbf{A}^{(1)}, \dots, \mathbf{A}^{(1)}] \\ \mathbf{B}(i) &= \begin{bmatrix} \mathbf{B}^{(1)}(j_{1i}) & & \\ & \ddots & \\ & & \mathbf{B}^{(1)}(j_{4i}) \end{bmatrix} \\ \mathbf{C} &= \mathbf{C}^{(1)} \\ \mathbf{D} &= \frac{1}{4} \begin{bmatrix} \mathbf{D}^{(1)} \\ \vdots \\ \mathbf{D}^{(1)} \end{bmatrix} \end{aligned} \quad (13)$$

Here, j_{ki} returns the index of the i -th control state under the k -th context; e.g., $j_{k\cdot} = [1, 2, 4, 3]$. This means that there are now 32 hidden states $S = S^{(2)} \otimes S^{(1)}$ (four

⁶For example, we do not have to worry about how the agent learns all possible configurations of the maze.

mazes, times four locations, times two reward contexts). The agent now has two levels of uncertainty to resolve. The first is induced by not knowing which maze it is in and the second is resolved by the cue, if it can be found. Neurobiologically, uncertainty about spatial context may be reflected in hippocampal processing (e.g., the “flickering” reported in Jezek, Henriksen, Treves, Moser, and Moser 2011).

Variational updating at the second level corresponds to replacing prior beliefs about hidden states with posterior beliefs after the previous trial. This corresponds to minimizing variational free energy because beliefs about the initial state $Q(s_0) = \mathbf{D} = \hat{s}_0$ become empirical priors that are informed by previous trials:

$$\begin{aligned}\hat{s}'_0 &= \mathbf{E}\hat{s}'_T \\ \mathbf{E} &= P(s'_0|s'_T, m) \\ &= ((1 - e)I_4 + e) \otimes (\mathbf{D}^{(1)} \otimes \mathbf{1}_8^T)\end{aligned}\quad (14)$$

Here, (s_0, s_T) and are the posterior expectations at the end of the previous trial and the beginning of the current trial respectively, and where $\mathbf{E} = P(s_0|s_T, m)$ encodes beliefs that the maze will change with a small probability $e = \frac{1}{8}$. This Bayesian belief updating is a formal way of saying that agents remember what they have learned from previous experience.

Figure 5, shows the results of this memory in terms of success rate as a function of trials, for eight successive trials in the same (randomly selected) mazes. The results are averaged over 128 realizations. Performance (gray lines) shows a progressive improvement as uncertainty about the hidden states falls (pink lines). This uncertainty is the entropy of posterior beliefs at the end of each trial $H[Q(s_T)]$ (multiplied by 100 for visual display). The equivalent performance for a conventional expected utility scheme is shown with dotted lines. The key thing to take from these results is that performance becomes near optimal after about four trials, at which point uncertainty falls to (nearly) zero. This means that, on average, the agent has learned which maze it is in after four trials and can then invoke the exploitative strategy of the previous section, first searching for the cue and then claiming the reward.

Crucially, despite the fact there is no explicit epistemic value involved in inference about the environment (maze) at the between-trial level, a failure to consider epistemic value at the within-trial level has deleterious consequences for learning, in that the expected utility agent fails to learn which

maze it is in (and is content to perform at the levels it would even if it knew). Note that the performance in Figure 5 never exceeds the performance shown in Figure 3.

The lower panels of Figure 5 show simulated dopamine responses (using the format of previous figures) over all iterations and trials. These results demonstrate the transfer of simulated dopamine responses to the cue or conditioned stimulus as learning progresses (gray arrows). These trial by trial changes are accompanied by elevated tonic responses after the CS—that reflect increasing confidence or precision about the outcomes of policies as the agent becomes familiar with its new environment (Hollerman & Schultz, 1996; Niv, 2007).

Summary

This section has shown it is straightforward to create hierarchical generative models, in which higher levels provide a context for lower levels, by equipping the model with a joint state-space $S = S^{(2)} \otimes S^{(1)}$ and associated transition matrices. This enables one to consider contingencies that are conserved (or not) over trials. In a multi-trial setting, priors over the initial state of each successive trial become empirical priors that minimize variational free energy (in exactly the same way as beliefs are updated within trials). This is simple to implement using Bayesian belief updating and allows a natural separation of temporal scales across hierarchical levels. It is relatively easy to see how one could generalize this to hierarchically deep models of the sort that real agents have to deal with, e.g., $S = S^{(3)} \otimes S^{(2)} \otimes S^{(1)}$.

This hierarchical augmentation reveals the role of integrating extrinsic and epistemic value in enabling the agent to learn which context it is operating in and then exploit that knowledge. It is tempting to associate this (inevitable and emergent) progression from exploration to exploitation with the transformation of goal-directed behavior into habits (Balleine & Dickinson, 1998; Dolan & Dayan, 2013; Pezzulo, Rigoli, & Chersi, 2013). Here, this Bayes optimal progression rests upon a contextualization of (first level) choice behavior by (second level) Bayesian updating that effectively accumulates evidence to resolve uncertainty about the consequences of behavior. This resolution restricts the repertoire of controlled state transitions that have to be considered in selecting the optimal policy and effectively increases the precision of action selection.

One might think that much of the delicate balance between exploration and exploitation could rest upon hierarchical active inference of this sort.

DISCUSSION

Formal approaches to decision-making under uncertainty generally rest on partially observable Markov decision processes, in which states are not directly observable but have to be inferred from observations. This formalism raises two fundamental issues that can be cast in terms of the exploration-exploitation dilemma. First, in relation to inference, in some circumstances an agent might obtain a larger reward by performing an epistemic (explorative) action rather than a more greedy (pragmatic) action. Second, in relation to learning, an epistemic action may be more appropriate to resolve uncertainty about aspects of its generative model. In classical formulations, the exploration-exploitation dilemma is usually solved with ad hoc solutions (like changing the precision of softmax decision rules). Here, we introduce a theoretical framework within which a solution to the exploration-exploitation dilemma emerges normatively from the minimization of expected free energy. For example, the precision or temperature parameter of softmax response rules becomes a parameter of the generative model and thereby acquires a Bayes optimal value.

More specifically, we have introduced a modeling framework for choice behavior that can be framed in terms of discrete states or (partially observed) Markov decision processes. There are two perspectives on this framework. People familiar with active inference could consider this work to show that the minimization of expected free energy furnishes a sufficient account of choice behavior under uncertainty. This necessarily entails epistemic action, providing a formal account of risk-sensitive or KL control and expected utility theory. The ensuing scheme also has construct validity in relation to Bayesian surprise and information theoretic formulations of search behavior. Crucially, the minimization of expected free energy eschews ad hoc parameters associated with conventional treatments (e.g., softmax parameters). Furthermore, active inference under hierarchical models may provide a useful framework within which to consider the contextualization of low-level behaviors that involves a natural (Bayes-optimal) progression from exploration to exploitation. Finally, it enables one to finesse the combinatorics of difficult or deep Markovian problems using approximate Bayesian inference—and a message passing scheme that is not biologically implausible. In

particular, the variational updates for expected precision show many similarities to empirical dopaminergic responses.

Our simulations suggest that it is difficult to completely suppress precision updates (dopaminergic responses), even when outcomes are very predictable (because every event portends something in our finite horizon setup). This contrasts with the classic results of Schultz and colleagues (Schultz et al., 1993), who found negligible responses to conditioned stimuli after learning. On the other hand, we were able to reproduce the empirical findings under conditions of uncertainty and predictive reward (Fiorillo et al., 2003; Schultz, 1998). Furthermore, the simulations reproduce the empirical observation that dopaminergic responses are transferred directly from the unconditioned stimuli to the conditioned stimuli, in the absence of any responses during the intervening period. Detailed response characteristics of this sort may provide important clues that may disambiguate, or further refine, theoretical accounts of dopaminergic function.

The second perspective on this work could be taken by people familiar with reinforcement learning, a branch of machine learning inspired by behavioral psychology (Sutton & Barto, 1998). From this perspective one can trace the steps that lead from normative descriptions based upon expected reward or utility to active inference and variational free energy minimization:

- The first step is to reformulate reinforcement learning or game theory problems as pure inference problems, i.e., planning as inference (Botvinick & Toussaint, 2012; Still, 2009; Still & Precup, 2012; Vijayakumar, Toussaint, Petkos, & Howard, 2009). This means that reward or utility functions become log probabilities defining prior beliefs or preferences about future outcomes. This induces probability distributions over policies that produce outcomes—and the precision of those distributions. This is important because it defines a Bayes-optimal precision for selecting among policies (Friston et al., 2014). Furthermore, casting reward or utility in terms of log probabilities means that they have the same currency as information (nats or bits), thereby providing a natural way to combine the value of an outcome and the value of information.
- The second step rests on accommodating uncertainty or risk over outcomes. When the expected utility of two choices is the same but one leads to several outcomes and the other a single outcome, then optimal behavior is not uniquely defined by expected utility. The simplest way to

accommodate uncertainty (risk) of this sort is to maximize both expected utility and the entropy of outcomes. Maximizing the expected entropy of outcomes effectively keeps one's options open (Klyubin, Polani, & Nehaniv, 2008). However, the sum of an entropy and expected utility can always be expressed as a (negative) KL divergence, leading to risk-sensitive or KL control. Formally, maximizing the expected utility and entropy of outcomes is equivalent to minimizing the KL divergence between the expected (predictive) distribution over outcomes and the distribution specified by the utility function. In behavioral terms, maximizing the entropy of controlled outcomes can be understood in terms of novelty bonuses and related concepts (Bach & Dolan, 2012; Daw et al., 2005; De Martino, Fleming, Garrett, & Dolan, 2012; Kakade & Dayan, 2002; Wittmann et al., 2008). In economics, there is a conceptual link with Shackle's formulation of *potential surprise* and the crucial role of money in facilitating risk-sensitive control (Shackle, 1972): If I am not sure what I want to buy, then I will save my money (liquid assets) and buy something later (maximize the entropy over future purchases—a fiscal Ockham's Razor).

- Risk-sensitive or KL control works fine if there is no uncertainty or ambiguity about hidden states given observed outcomes. However, when the same state can lead to several outcomes (e.g., noisy or ambiguous cues), we have to augment the KL divergence with the expected entropy over outcomes given the hidden states that cause them. Minimizing this entropy ensures that hidden states generating ambiguous (high entropy) outcomes are avoided. In other words, observations that resolve uncertainty about hidden states become intrinsically valuable. However, the sum of the expected conditional entropy and the KL divergence is the expected free energy that scores the quality or value of a policy. This brings us to active inference and the minimization of expected free energy that is sensitive to both risk and ambiguity.

In what follows, we consider some of the theoretical implications of these arguments, in relation to established approaches in psychology and artificial intelligence.

Curiosity and Bayesian surprise

Epistemic value and implicit exploratory behavior are related to curiosity in psychology (Harlow, 1950;

Ryan & Deci, 1985) and intrinsic motivation in reinforcement learning (Baldassarre & Mirolli, 2013; Barto, Singh, & Chentanez, 2004; Oudeyer & Kaplan, 2007; Schembri, Mirolli, & Baldassarre, 2007; Schmidhuber, 1991). Here *intrinsic* stands in opposition to *extrinsic* (e.g., drive or goal) value. While we have focused on reducing uncertainty during inference, most reinforcement learning research uses curiosity or novelty-based mechanisms to learn a policy or model efficiently. The general idea here is that an agent should select actions that improve learning or prediction, thus avoiding behaviors that preclude learning (either because these behaviors are already learned or because they are unlearnable). It has often been emphasized that adaptive agents should seek out *surprising* stimuli, not *unsurprising* stimuli as assumed in active inference. This apparent discrepancy can be reconciled if one considers that surprising events, in the setting of curiosity and Bayesian surprise, are simply outcomes that are *salient* and minimize uncertainty. In active inference, agents are surprised when they do not minimize uncertainty. It is salient (counterfactual) outcomes that optimize exploration (and model selection) and salience-seeking behavior stems nicely from the more general objective of minimizing expected free energy (or surprise proper).

There is, however, an important difference between active inference and the concepts of curiosity and Bayesian surprise, at least as they are usually used. Salience is typically framed in “bottom-up” terms, in that the agents are not assumed to have a particular goal or task. This is also a characteristic of curiosity (and similar) algorithms that try to learn all possible models, without knowing in advance which will be useful for achieving a specific goal. The active inference scheme considered here contextualizes the utilitarian value of competing policies in terms of their epistemic value, where the implicit reduction in uncertainty is (or can be) tailored for the goals or preferred outcomes in mind.

Active inference and the exploitation-exploration dilemma

The active inference formulation effectively combines belief state updates, action selection, and learning under a single imperative. In principle, this results in the efficient learning of both the structure of the environment and the selection of the suitable policies, thereby avoiding the problems of model-free reinforcement learning algorithms (Sutton & Barto, 1998). Model-free schemes need to relearn a policy

every time the environment changes (Ognibene, Pezzulo, & Baldassarre, 2010). Active inference offers a principled solution to the exploration-exploitation dilemma and, in contrast with model-based learning, will not waste time modeling irrelevant aspects of the environment (Atkeson & Santamaria, 1997). This may enhance learning through generalization, by predominantly sampling features that are conserved when the environmental context changes (Ognibene & Baldassarre, 2014; Walther, Rutishauser, Koch, & Perona, 2005).

Furthermore, active inference extends established metaphors for purely perceptual processing, in particular, hierarchical Bayesian filtering and predictive coding (Clark, 2013; Friston, 2010; Lee & Mumford, 2003; Rao & Ballard, 1999). These perspectives can explain several aspects of cortical hierarchies (Dayan et al., 1995) and provide a nice perspective on the brain as an organ that adapts to model and predict its sensory inputs. This is particularly important because the resulting hierarchical representation (deep generative model) can account for sensorimotor regularities produced by action (Lungarella & Sporns, 2006; O'Regan & Noë, 2001). In turn, this can improve learning and inference, which depend sensitively on an efficient and sparse (hierarchical) representation of active sampling and sensorimotor learning (Ballard et al., 2013; Barto et al., 2004; Tani & Nolfi, 1999). From a modeling perspective, the integration of learning, belief updating, and action selection may allow one to study, in a principled manner, how perception supports learning and how learning can result in different internal representations (Little & Sommer, 2013; Lungarella & Sporns, 2006; Ognibene & Baldassarre, 2014; Verschure, Voegtlin, & Douglas, 2003).

This may be particularly important when modeling inference and behavior in tasks where the agent has no detailed knowledge of the environment, e.g., foraging in open and changing environments, possibly with other agents. These difficult problems have limited the application of the MDP framework to tasks with definitive and detailed representations, such as navigation in grid-based mazes. In open environments, epistemic behaviors have been largely described with heuristic (Brooks, 1991; Itti & Koch, 2001) or stochastic processes such as Lévy flight (Beer, 1995; Viswanathan et al., 1999). However, modeling these problems within the active inference framework may reveal the formal nature of these processes and their neuronal correlates.

Bayesian Reinforcement Learning (e.g., Cao & Ray, 2012) also provides a principled approach to the exploration-exploitation trade-off and explicitly models uncertainty about the quality of alternative policies. Because active inference tackles both the

problems of learning and of exploration under partial observations in a coherent manner, it would be interesting to see if Bayesian reinforcement learning could be formulated in terms of active inference. This may be useful, because the current scheme offers computational efficiency, by exploiting variational Bayesian techniques (c.f., Friston & Barber, 2010), accommodates formal constraints on the structure of policies, and comes with a biologically plausible process theory.

Applications

While these are clearly interesting theoretical issues, the purpose of this paper is also pragmatic. The simulations presented in this paper all use one (Matlab) routine that only requires the specification of the generative model in terms of its $(\mathbf{A}, \mathbf{B}, \mathbf{C}, \mathbf{D}) \subset \theta$ parameters. Crucially, integrating this scheme, for any given set of choices and outcomes, provides a generative model of empirical choice behavior. This means, one can estimate the parameters that are unique to a particular subject (human or animal) using standard (meta-Bayesian) schemes (Daunizeau et al., 2010). These parameters include the sensitivity to particular outcomes, beliefs about experimental contingencies, and the overall confidence (and confidence in confidence) encoded by a subject's hyperpriors over precision, e.g., $(c, a, \alpha, \beta) \subset \theta$. This enables a cognitive and possibly physiological phenotyping of subjects using behavioral and physiological responses respectively. Furthermore, one could use Bayesian model comparison to assess whether subjects use expected utility, risk-sensitive control, or full active inference. Indeed, we have shown that the choice behavior and fMRI responses in the dopaminergic midbrain area are better explained in terms of KL control, relative to expected utility using this approach (Schwartenbeck, FitzGerald, Mathys, Dolan, & Friston, 2014). We hope to pursue a similar approach to exploration and decision-making under uncertainty in future work.

Original manuscript received 23 October 2014

Revised manuscript received 25 January 2015

First published online 13 March 2015

REFERENCES

Andreopoulos, A., & Tsotsos, J. (2013). A computational learning theory of active object recognition under

- uncertainty. *International Journal of Computer Vision*, 101(1), 95–142. doi:[10.1007/s11263-012-0551-6](https://doi.org/10.1007/s11263-012-0551-6)
- Atkeson, C., & Santamaría, J. (1997). A comparison of direct and model-based reinforcement learning. In *International Conference on Robotics and Automation (ICRA)*. IEEE Press.
- Attias, H. (2003). Planning by probabilistic inference. In C. M. Bishop and B. J. Frey (Eds.), *Proceedings of the Ninth International Workshop on Artificial Intelligence and Statistics*. Key West, FL.
- Bach, D. R., & Dolan, R. J. (2012). Knowing how much you don't know: A neural organization of uncertainty estimates. *Nature Reviews. Neuroscience*, 13(8), 572–586.
- Baldassarre, G., & Mirolli, M. (2013). *Intrinsically motivated learning in natural and artificial systems*. Springer: Berlin.
- Ballard, D. H., Kit, D., Rothkopf, C. A., & Sullivan, B. (2013). A hierarchical modular architecture for embodied cognition. *Multisensory Research*, 26, 177. doi:[10.1163/22134808-00002414](https://doi.org/10.1163/22134808-00002414)
- Balleine, B. W., & Dickinson, A. (1998). Goal-directed instrumental action: Contingency and incentive learning and their cortical substrates. *Neuropharmacology*, 37 (4–5), 407–419. doi:[10.1016/S0028-3908\(98\)00033-1](https://doi.org/10.1016/S0028-3908(98)00033-1)
- Barlow, H. (1961). Possible principles underlying the transformations of sensory messages. In W. Rosenblith (Ed.), *Sensory communication* (pp. 217–234). Cambridge, MA: MIT Press.
- Barlow, H. B. (1974). Inductive inference, coding, perception, and language. *Perception*, 3, 123–134. doi:[10.1068/p030123](https://doi.org/10.1068/p030123)
- Barto, A., Singh, S., & Chentanez, N. (2004). Intrinsically motivated learning of hierarchical collections of skills. In *Proceedings of the 3rd International Conference on Development and Learning (ICDL 2004)*, San Diego: Salk Institute.
- Beal, M. J. (2003). *Variational algorithms for approximate bayesian inference* (PhD. Thesis). University College London.
- Beer, R. D. (1995). A dynamical systems perspective on agent-environment interaction. *Artificial Intelligence*, 72 (1–2), 173–215. doi:[10.1016/0004-3702\(94\)00005-L](https://doi.org/10.1016/0004-3702(94)00005-L)
- Berridge, K. C. (2007). The debate over dopamine's role in reward: The case for incentive salience. *Psychopharmacology (Berl.)*, 191(3), 391–431. doi:[10.1007/s00213-006-0578-x](https://doi.org/10.1007/s00213-006-0578-x)
- Bialek, W., Nemenman, I., & Tishby, N. (2001). Predictability, complexity, and learning. *Neural Computation*, 13(11), 2409–2463. doi:[10.1080/01621459.1989.10478825](https://doi.org/10.1080/01621459.1989.10478825)
- Bonet, B., & Geffner, H. (2014). Belief tracking for planning with sensing: Width, complexity and approximations. *Journal of Artificial Intelligence Research*, 50, 923–970.
- Botvinick, M., & An, J. (2008). Goal-directed decision making in prefrontal cortex: A computational framework. *Advances in Neural Information Processing Systems (NIPS)*.
- Botvinick, M., & Toussaint, M. (2012). Planning as inference. *Trends in Cognitive Sciences*, 16(10), 485–488. doi:[10.1016/j.tics.2012.08.006](https://doi.org/10.1016/j.tics.2012.08.006)
- Braun, D. A., Ortega, P. A., Theodorou, E., & Schaal, S. (2011). Path integral control and bounded rationality. *Adaptive Dynamic Programming And Reinforcement Learning (ADPRL)*, 2011 IEEE Symposium on, Paris, IEEE.
- Bromberg-Martin, E. S., & Hikosaka, O. (2009). Midbrain dopamine neurons signal preference for advance information about upcoming rewards. *Neuron*, 63(1), 119–126. doi:[10.1016/j.neuron.2009.06.009](https://doi.org/10.1016/j.neuron.2009.06.009)
- Brooks, R. (1991). Intelligence without representation. *Artificial Intelligence*, 47(1–3), 139–159. doi:[10.1016/0004-3702\(91\)90053-M](https://doi.org/10.1016/0004-3702(91)90053-M)
- Bruce, N., & Tsotsos, J. (2009). Saliency, attention, and visual search: An information theoretic approach. *Journal of Vision*, 9 (3):5, 1–24. doi:[10.1167/9.3.5](https://doi.org/10.1167/9.3.5)
- Bunzeck, N., & Düzel, E. (2006). Absolute coding of stimulus novelty in the human substantia nigra/VTa. *Neuron*, 51(3), 369–379. doi:[10.1016/j.neuron.2006.06.021](https://doi.org/10.1016/j.neuron.2006.06.021)
- Canolty, R. T., Edwards, E., Dalal, S. S., Soltani, M., Nagarajan, S. S., Kirsch, H. E. ... Knight, R. T. (2006). High gamma power is phase-locked to theta oscillations in human neocortex. *Science*, 313(5793), 1626–1628. doi:[10.1126/science.1128115](https://doi.org/10.1126/science.1128115)
- Cao, F., & Ray, S. (2012). Bayesian hierarchical reinforcement learning. *Advances in Neural Information Processing Systems*.
- Cisek, P. (2007). Cortical mechanisms of action selection: The affordance competition hypothesis. *Philosophical Transactions of the Royal Society of London. Series B, Biological Sciences*, 362(1485), 1585–1599.
- Clark, A. (2013). Whatever next? Predictive brains, situated agents, and the future of cognitive science. *The Behavioral and Brain Sciences*, 36(3), 181–204. doi:[10.1017/S0140525X12000477](https://doi.org/10.1017/S0140525X12000477)
- Cohen, J. D., McClure, S. M., & Yu, A. J. (2007). Should I stay or should I go? How the human brain manages the trade-off between exploitation and exploration. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 362(1481), 933–942. doi:[10.1098/rstb.2007.2098](https://doi.org/10.1098/rstb.2007.2098)
- D'Ardenne, K., McClure, S. M., Nystrom, L. E., & Cohen, J. D. (2008). BOLD responses reflecting dopaminergic signals in the human ventral tegmental area. *Science*, 319(5867), 1264–1267. doi:[10.1126/science.1150605](https://doi.org/10.1126/science.1150605)
- Daunizeau, J., Den Ouden, H. E., Pessiglione, M., Kiebel, S. J., Stephan, K. E., & Friston, K. J. (2010). Observing the observer (I): Meta-bayesian models of learning and decision-making. *PLoS One*, 5(12), e15554. doi:[10.1371/journal.pone.0015554](https://doi.org/10.1371/journal.pone.0015554)
- Daw, N. D., & Doya, K. (2006). The computational neurobiology of learning and reward. *Current Opinion in Neurobiology*, 16(2), 199–204. doi:[10.1016/j.conb.2006.03.006](https://doi.org/10.1016/j.conb.2006.03.006)
- Daw, N. D., Niv, Y., & Dayan, P. (2005). Uncertainty-based competition between prefrontal and dorsolateral striatal systems for behavioral control. *Nature Neuroscience*, 8 (12), 1704–1711. doi:[10.1038/nn1560](https://doi.org/10.1038/nn1560)
- Daw, N. D., O'Doherty, J. P., Dayan, P., Seymour, B., & Dolan, R. J. (2006). Cortical substrates for exploratory decisions in humans. *Nature*, 441(7095), 876–879. doi:[10.1038/nature04766](https://doi.org/10.1038/nature04766)
- Dayan, P. (2009). Dopamine, reinforcement learning, and addiction. *Pharmacopsychiatry*, 42(1), S56–65. doi:[10.1055/s-0028-1124107](https://doi.org/10.1055/s-0028-1124107)
- Dayan, P., & Hinton, G. E. (1997). Using expectation maximization for reinforcement learning. *Neural*

- Computation*, 9, 271–278. doi:10.1162/neco.1997.9.2.271
- Dayan, P., Hinton, G. E., & Neal, R. (1995). The Helmholtz machine. *Neural Computation*, 7, 889–904. doi:10.1162/neco.1995.7.5.889
- De Martino, B., Fleming, S. M., Garrett, N., & Dolan, R. J. (2012). Confidence in value-based choice. *Nature Neuroscience*, 16, 105–110. doi:10.1038/nn.3279
- Dolan, R. J., & Dayan, P. (2013). Goals and habits in the brain. *Neuron*, 80(2), 312–325. doi:10.1016/j.neuron.2013.09.007
- Ferro, M., Ognibene, D., Pezzulo, G., & Pirrelli, V. (2010). Reading as active sensing: A computational model of gaze planning during word recognition. *Frontiers in Neurobotics*, 4, 1.
- Fiorillo, C. D., Tobler, P. N., & Schultz, W. (2003). Discrete coding of reward probability and uncertainty by dopamine neurons. *Science*, 299(5614), 1898–1902. doi:10.1126/science.1077349
- FitzGerald, T., Dolan, R., & Friston, K. (2014). Model averaging, optimal inference, and habit formation. *Frontiers in Human Neuroscience*. doi:10.3389/fnhum.2014.00457
- Fox, C., & Roberts, S. (2011). A tutorial on variational Bayes. *Artificial Intelligence Review*, Springer. doi:10.1007/s10462-01011-19236-10468
- Friston, K. (2010). The free-energy principle: A unified brain theory? *Nature Reviews Neuroscience*, 11(2), 127–138. doi:10.1038/nrn2787
- Friston, K., Adams, R. A., Perrinet, L., & Breakspear, M. (2012). Perceptions as hypotheses: Saccades as experiments. *Frontiers in Psychology*, 3, 151. doi:10.3389/fpsyg.2012.00151
- Friston, K., & Mathys, C. (2015). I think therefore I am. Cognitive Dynamic Systems. S. Haykin, IEEE press: in press.
- Friston, K., Samothrakis, S., & Montague, R. (2012). Active inference and agency: Optimal control without cost functions. *Biological Cybernetics*, 106, 523–541. [Epub ahead of print]. doi:10.1007/s00422-012-0512-8
- Friston, K., Schwartenbeck, P., FitzGerald, T., Moutoussis, M., Behrens, T., & Dolan, R. J. (2014). The anatomy of choice: Dopamine and decision-making. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 369(1655). doi:10.1098/rstb.2013.0481
- Friston, K., Schwartenbeck, P., FitzGerald, T., Moutoussis, M., Behrens, T., Raymond, R. J., & Dolan, J. (2013). The anatomy of choice: Active inference and agency. *Frontiers in Human Neuroscience*, 7, 598. doi:10.3389/fnhum.2013.00598
- Furmston, T., & Barber, D. (2010). Variational methods for reinforcement learning. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*. JMLR: W&CP, vol. 9, (pp. 241–248).
- Fuster, J. M. (2004). Upper processing stages of the perception-action cycle. *Trends in Cognitive Sciences*, 8(4), 143–145. doi:10.1016/j.tics.2004.02.004
- Gurney, K., Prescott, T. J., & Redgrave, P. (2001). A computational model of action selection in the basal ganglia. I. A new functional anatomy. *Biological Cybernetics*, 84(6), 401–410. doi:10.1007/PL00007984
- Harlow, H. (1950). Learning and satiation of response to intrinsically motivated complex puzzle performance by monkeys.” *Journal of Comparative Physiological Psychology*, 40, 289–294. doi:10.1037/h0058114
- Hauskrecht, M. (2000). Value-function approximations for partially observable Markov decision processes. *Journal of Artificial Intelligence Research*, 13, 33–94.
- Helmholtz, H. (1866/1962). *Concerning the perceptions in general. Treatise on physiological optics*. New York, NY: Dover. III.
- Hollerman, J. R., & Schultz, W. (1996). Activity of dopamine neurons during learning in a familiar task context. *Social Neuroscience Abstracts*, 22, 1388.
- Howard, R. (1966). Information value theory. *IEEE Transactions on Systems, Science and Cybernetics*, SSC-2 (1), 22–26. doi:10.1109/TSSC.1966.300074
- Humphries, M. D., Khamassi, M., & Gurney, K. (2012). Dopaminergic control of the exploration-exploitation trade-off via the basal ganglia. *Frontiers in Neuroscience*, 6, 9. doi:10.3389/fnins.2012.00009
- Itti, L., & Baldi, P. (2009). Bayesian surprise attracts human attention. *Vision Research*, 49(10), 1295–1306. doi:10.1016/j.visres.2008.09.007
- Itti, L., & Koch, C. (2001). Computational modelling of visual attention. *Nature Reviews Neuroscience*, 2(3), 194–203. doi:10.1038/35058500
- Jezek, K., Henriksen, E., Treves, A., Moser, E., & Moser, M.-B. (2011). Theta-paced flickering between place-cell maps in the hippocampus. *Nature*, 478, 246–249. doi:10.1038/nature10439
- Kaelbling, L. P., Littman, M. L., & Cassandra, A. R. (1998). Planning and acting in partially observable stochastic domains. *Artificial Intelligence*, 101(1–2), 99–134. doi:10.1016/S0004-3702(98)00023-X
- Kakade, S., & Dayan, P. (2002). Dopamine: Generalization and bonuses. *Neural Networks*, 15(4–6), 549–559. doi:10.1016/S0893-6080(02)00048-5
- Kamar, E., & Horvitz, E. (2013). Light at the end of the tunnel: A Monte Carlo approach to computing value of information. In *Proceedings of the 2013 international conference on Autonomous agents and multi-agent systems*.
- Kappen, H. J., Gomez, Y., & Opper, M. (2012). Optimal control as a graphical model inference problem. *Machine Learning*, 87(2), 159–182. doi:10.1007/s10994-012-5278-7
- Kass, R. E., & Steffey, D. (1989). Approximate Bayesian inference in conditionally independent hierarchical models (parametric empirical Bayes models). *Journal of the American Statistical Association*, 407, 717–726.
- Kirsh, D., & Maglio, P. (1994). On distinguishing epistemic from pragmatic action. *Cognitive Science*, 18(4), 513–549. doi:10.1207/s15516709cog1804_1
- Klyubin, A. S., Polani, D., & Nehaniv, C. L. (2008). Keep your options open: An information-based driving principle for sensorimotor systems. *PLoS One*, 3(12), e4018. doi:10.1371/journal.pone.0004018
- Krause, A., & Guestrin, C. (2005). Optimal nonmyopic value of information in graphical models: Efficient algorithms and theoretical limits. In *International Joint Conference on Artificial Intelligence (IJCAI)*. Carnegie Mellon University.
- Krebs, R. M., Schott, B. H., Schütze, H., & Düzel, E. (2009). The novelty exploration bonus and its

- attentional modulation. *Neuropsychologia*, 47, 2272–2281. doi:10.1016/j.neuropsychologia.2009.01.015
- Lee, T. S., & Mumford, D. (2003). Hierarchical Bayesian inference in the visual cortex. *Journal of the Optical Society of America A*, 20, 1434–1448. doi:10.1364/JOSAA.20.001434
- Lepora, N., Martinez-Hernandez, U., & Prescott, T. (2013). Active touch for robust perception under position uncertainty. In *IEEE proceedings of ICRA*.
- Linsker, R. (1990). Perceptual neural organization: Some approaches based on network models and information theory. *Annual Review of Neuroscience*, 13, 257–281. doi:10.1146/annurev.ne.13.030190.001353
- Little, D. Y., & Sommer, F. T. (2013). Learning and exploration in action-perception loops. *Front Neural Circuits*, 7, 37. doi:10.3389/fncir.2013.00037
- Littman, M., Sutton, R., & Singh, S. (2002). Predictive Representations of State. *Advances in Neural Information Processing Systems 14 (NIPS)*.
- Lungarella, M., & Sporns, O. (2006). Mapping information flow in sensorimotor networks. *PLoS Computational Biology*, 2, e144. doi:10.1371/journal.pcbi.0020144
- McClure, S. M., Daw, N. D., & Montague, P. R. (2003). A computational substrate for incentive salience. *Trends in Neurosciences*, 26(8), 423–428. doi:10.1016/S0166-2236(03)00177-2
- Miller, G., Galanter, E., & Pribram, K. (1960). *Plans and the structure of behavior*. New York, NY: Henry Holt.
- Moser, E. I., Kropff, E., & Moser, M.-B. (2008). Place cells, grid cells, and the brain's spatial representation system. *Annual Review of Neuroscience*, 31, 69–89. doi:10.1146/annurev.neuro.31.061307.090723
- Mushiake, H., Saito, N., Sakamoto, K., Itoyama, Y., & Tanji, J. (2006). Activity in the lateral prefrontal cortex reflects multiple steps of future events in action plans. *Neuron*, 50, 631–641. doi:10.1016/j.neuron.2006.03.045
- Najemnik, J., & Geisler, W. (2005). Optimal eye movement strategies in visual search. *Nature*, 434, 387–391. doi:10.1038/nature03390
- Nepora, N., & Gurney, K. (2012). The basal ganglia optimize decision making over general perceptual hypotheses. *Neural Computation*, 24(11), 2924–2945. doi:10.1162/NECO_a_00360
- Niv, Y. (2007). Cost, benefit, tonic, phasic: What do response rates tell us about dopamine and motivation? *Annals of the New York Academy of Sciences*, 1104, 357–376.
- O'Doherty, J., Dayan, P., Schultz, J., Deichmann, R., Friston, K., & Dolan, R. J. (2004). Dissociable roles of ventral and dorsal striatum in instrumental conditioning. *Science*, 304(5669), 452–454. doi:10.1126/science.1094285
- O'Regan, J., & Noë, A. (2001). A sensorimotor account of vision and visual consciousness. *The Behavioral and Brain Sciences*, 24, 939–973. doi:10.1017/S0140525X01000115
- Ognibene, D., & Baldassarre, G. (2014). Ecological active vision: Four Bio-inspired principles to integrate bottom-up and adaptive top-down attention tested with a simple camera-arm robot. *Autonomous Mental Development, IEEE Transactions on*. IEEE.
- Ognibene, D., Chinellato, E., Sarabia, M., & Demiris, Y. (2013). Contextual action recognition and target localization with an active allocation of attention on a humanoid robot. *Bioinspiration & Biomimetics*, 8, 3. doi:10.1088/1748-3182/8/3/035002
- Ognibene, D., & Demiris, Y. (2013). Toward active event perception. In *International Joint conference of Artificial Intelligence*. IJCAI: Beijing.
- Ognibene, D., Pezzulo, G., & Baldassarre, G. (2010). Learning to look in different environments: An active-vision model which learns and readapts visual routines. In S. Doncieux (ed.), *Proceedings of the 11th International Conference on Simulation of Adaptive Behaviour*, SAB 2010, Paris - Clos Lucé, France, August 25–28, 2010. *Proceedings*. Springer Berlin Heidelberg.
- Ognibene, D., Catenacci Volpi, N., Pezzulo, G., & Baldassarre, G. (2013). Learning epistemic actions in model-free memory-free reinforcement learning: Experiments with a neuro-robotic model. In *Second International Conference on Biomimetic and Biohybrid Systems*. *Proceedings*. pp. 191–203. doi:10.1007/978-3-642-39802-5_17
- Oja, E. (1989). Neural networks, principal components, and subspaces. *International Journal of Neural Systems*, 1, 61–68. doi:10.1142/S0129065789000475
- Olshausen, B. A., & Field, D. J. (1996). Emergence of simple-cell receptive field properties by learning a sparse code for natural images. *Nature*, 381, 607–609. doi:10.1038/381607a0
- Optican, L., & Richmond, B. J. (1987). Temporal encoding of two-dimensional patterns by single units in primate inferior cortex. II Information Theoretic Analysis. *Journal of Neurophysiology*, 57, 132–146.
- Ortega, P. A., & Braun, D. A. (2011). Information, utility and bounded rationality. *Lecture Notes on Artificial Intelligence*, 6830, 269–274.
- Ortega, P. A., & Braun, D. A. (2013). Thermodynamics as a theory of decision-making with information-processing costs. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 469, 2153. doi:10.1098/rspa.2012.0683
- Oudeyer, P.-Y., & Kaplan, F. (2007). What is intrinsic motivation? A typology of computational approaches. *Frontiers in Neurobotics*, 1, 6. doi:10.3389/neuro.12.006.2007
- Penny, W., Zeidman, P., & Burgess, N. (2013). Forward and backward inference in spatial cognition. *Plos Computational Biology*, 9(12), e1003383. doi:10.1371/journal.pcbi.1003383
- Pessiglione, M., Seymour, B., Flandin, G., Dolan, R. J., & Frith, C. D. (2006). Dopamine-dependent prediction errors underpin reward-seeking behaviour in humans. *Nature*, 442(7106), 1042–1045. doi:10.1038/nature05051
- Pezzementi, Z., Plaku, E., Reyda, C., & Hager, G. (2011). Tactile-object recognition from appearance information. *IEEE Transactions on Robotics*, 27, 473–487. doi:10.1109/TRO.2011.2125350
- Pezzulo, G., & Castelfranchi, C. (2009). Thinking as the control of imagination: A conceptual framework for goal-directed systems. *Psychological Research Psychologische Forschung*, 73, 559–577. doi:10.1007/s00426-009-0237-z
- Pezzulo, G., Rigoli, F., & Chersi, F. (2013). The mixed instrumental controller: Using value of information to combine habitual choice and mental simulation.

- Frontiers in Psychology*, 4, 92. doi:[10.3389/fpsyg.2013.00092](https://doi.org/10.3389/fpsyg.2013.00092)
- Pezzulo, G., Van Der Meer, M., Lansink, C., & Pennartz, C. (2014). Internally generated sequences in learning and executing goal-directed behavior. *Trends in Cognitive Sciences*, 18(2), pp. 647–657.
- Rao, R. P., & Ballard, D. H. (1999). Predictive coding in the visual cortex: A functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2(1), 79–87. doi:[10.1038/4580](https://doi.org/10.1038/4580)
- Redgrave, P., & Gurney, K. (2006). The short-latency dopamine signal: A role in discovering novel actions? *Nature Reviews. Neuroscience*, 7(12), 967–975. doi:[10.1038/nrn2022](https://doi.org/10.1038/nrn2022)
- Roy, N., Burgard, W., Fox, D., & Thrun, S. (1999). Coastal navigation: Robot navigation under uncertainty in dynamic environments. In *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*. Washington, DC: IEEE Computer Society.
- Ryan, R., & Deci, E. (1985). *Intrinsic motivation and self-determination in human behavior*. New York, NY: Plenum.
- Schembri, M., Mirolli, M., & Baldassare, G. (2007). Evolving internal reinforcers for an intrinsically motivated reinforcement learning robot. In Y. Demiris, B. Scassellati, & D. Mareschal (Eds.) *Proceedings of the 6th IEEE International Conference on Development and Learning (ICDL2007)*. IEEE.
- Schmidhuber, J. (1991). Curious model-building control systems. In *Proceeding of International Joint Conference on Neural Networks*. Singapore: IEEE. 2: 1458–1463.
- Schneider, A., Sturm, J., Stachniss, C., Reiser, M., Burkhardt, H., & Burgard, W. (2009). Object identification with tactile sensors using bag-of-features. *IROS 2009 IEEE/RSJ International Conference on Intelligent Robots and Systems, IROS 2009*. p. 243. IEEE.
- Schultz, W. (1998). Predictive reward signal of dopamine neurons. *Journal of Neurophysiology*, 80(1), 1–27.
- Schultz, W., Apicella, P., & Ljungberg, T. (1993). Responses of monkey dopamine neurons to reward and conditioned stimuli during successive steps of learning a delayed response task. *Journal of Neuroscience*, 13, 900–913.
- Schwartenbeck, P., FitzGerald, T., Mathys, C., Dolan, R., & Friston, K. (2014). The dopaminergic midbrain encodes the expected certainty about desired outcomes. *Cerebral Cortex*. pii: bhu159. doi:[10.1093/cercor/bhu159](https://doi.org/10.1093/cercor/bhu159)
- Shackle, G. (1972). *Epistemics and economics*. Cambridge, MA: Cambridge University Press.
- Singh, A., Krause, A., Guestrin, C., & Kaiser, W. (2009). Efficient informative sensing using multiple robots. *Journal of Artificial Intelligence Research*, 34(2), 707.
- Solway, A., & Botvinick, M. (2012). Goal-directed decision making as probabilistic inference: A computational framework and potential neural correlates. *Psychological Review*, 119, 120–154. doi:[10.1037/a0026435](https://doi.org/10.1037/a0026435)
- Sornkarn, N., Nanayakkara, T., & Howard, M. (2014). Internal impedance control helps information gain in embodied perception. In *IEEE International Conference on Robotics and Automation (ICRA)*, IEEE.
- Still, S. (2009). Information-theoretic approach to interactive learning. *EPL (Europhysics Letters)*, 85(2), 28005. doi:[10.1209/0295-5075/85/28005](https://doi.org/10.1209/0295-5075/85/28005)
- Still, S., & Precup, D. (2012). An information-theoretic approach to curiosity-driven reinforcement learning. *Theory in Biosciences*, 131(3), 139–148. doi:[10.1007/s12064-011-0142-z](https://doi.org/10.1007/s12064-011-0142-z)
- Sutton, R. S., & Barto, A. G. (1998). *Reinforcement learning: An introduction*. Cambridge, MA: MIT Press.
- Tani, J., & Nolfi, S. (1999). Learning to perceive the world as articulated: An approach for hierarchical learning in sensory-motor systems. *Neural Networks*, 12, 1131–1141. doi:[10.1016/S0893-6080\(99\)00060-X](https://doi.org/10.1016/S0893-6080(99)00060-X)
- Tishby, N., & Polani, D. (2010). Information theory of decisions and actions. In V. Cutsuridis, A. Hussain, & J. Taylor (Eds.), *Perception-reason-action cycle: Models, algorithms and systems*. Springer: Berlin.
- Toussaint, M., & Storkey, A. (2006). Probabilistic inference for solving discrete and continuous state Markov Decision Processes. In *Proceeding of the 23rd International Conference on Machine Learning*. ACM. pp. 945–952.
- Van den Broek, J. L., Wiegerinck, W. A. J. J., & Kappen, H. J. (2010). Risk-sensitive path integral control. *UAI*, 6, 1–8.
- Verschure, P. F., Voegtlin, T., & Douglas, R. J. (2003). Environmentally mediated synergy between perception and behavior in mobile robots. *Nature*, 425, 620–624. doi:[10.1038/nature02024](https://doi.org/10.1038/nature02024)
- Vijayakumar, S., Toussaint, M., Petkos, G., & Howard, M. (2009). Planning and moving in dynamic environments. In B. Sendhoff (Ed.), *Creating brain-like intelligence* (pp. 151–191). Berlin: Springer-Verlag.
- Viswanathan, G., Buldyrev, S., Havlin, S., Da Luz, M., Raposo, E., & Stanley, H. (1999). Optimizing the success of random searches. *Nature*, 401(6756), 911–914. doi:[10.1038/44831](https://doi.org/10.1038/44831)
- Walther, D., Rutishauser, U., Koch, C., & Perona, P. (2005). Selective visual attention enables learning and recognition of multiple objects in cluttered scenes. *Computer Vision and Image Understanding*, 100, 41–63. doi:[10.1016/j.cviu.2004.09.004](https://doi.org/10.1016/j.cviu.2004.09.004)
- Wittmann, B. C., Daw, N. D., Seymour, B., & Dolan, R. J. (2008). Striatal activity underlies novelty-based choice in humans. *Neuron*, 58(6), 967–973. doi:[10.1016/j.neuron.2008.04.027](https://doi.org/10.1016/j.neuron.2008.04.027)

APPENDIX A

Lemma (predictive free energy): Under a generative model $P(s_\tau, o_\tau | \pi) = Q(s_\tau | o_\tau, \pi)P(o_\tau | m)$ and policy π , the negative free energy of the approximate posterior predictive density is $\forall : \tau > t$

$$\mathbf{Q}_\tau(\pi) = - \underbrace{E_{Q(s_\tau | \pi)}[H[P(o_\tau | s_\tau)]]}_{\text{Predicted uncertainty}} - \underbrace{D[Q(o_\tau | \pi) || P(o_\tau | m)]}_{\text{Predicted divergence}} \quad \text{A1.1}$$

Proof: The expected free energy of the approximate posterior predictive distribution over hidden states (under policy π at $\tau > t$ in the future) is the expected energy minus its entropy (where the energy of a hidden state $G(s_\tau, \pi)$ is itself an expectation over outcomes):

$$\begin{aligned} G(s_\tau, \pi) &= -E_{P(o_\tau | s_\tau)}[\ln P(o_\tau, s_\tau | \pi)] \\ F_\tau(\pi) &= E_{Q(s_\tau | \pi)}[G(s_\tau, \pi)] - H[Q(s_\tau | \pi)] \end{aligned} \quad \text{A1.2}$$

This means the quality or value of the policy is:

$$\begin{aligned} \mathbf{Q}_\tau(\pi) &= -F_\tau(\pi) \\ &= E_{Q(o_\tau, s_\tau | \pi)}[\ln P(o_\tau, s_\tau | \pi) - \ln Q(s_\tau | \pi)] \\ &= E_{Q(o_\tau, s_\tau | \pi)}[\ln Q(s_\tau | o_\tau, \pi) + \ln P(o_\tau | m) - \ln Q(s_\tau | \pi)] \\ &= E_{Q(o_\tau, s_\tau | \pi)}[\ln Q(o_\tau | s_\tau, \pi) + \ln P(o_\tau | m) - \ln Q(o_\tau | \pi)] \\ &= - \underbrace{E_{Q(s_\tau | \pi)}[H[P(o_\tau | s_\tau)]]}_{\text{Predicted uncertainty}} - \underbrace{D[Q(o_\tau | \pi) || P(o_\tau | m)]}_{\text{Predicted divergence}} \end{aligned}$$

$$\begin{aligned} Q(s_\tau | \pi) &= E_{Q(s_\tau)}[P(s_\tau | s_t, \pi)] \\ Q(o_\tau | \pi) &= E_{Q(s_\tau | \pi)}[P(o_\tau | s_\tau)] \\ Q(o_\tau, s_\tau | \pi) &= P(o_\tau | s_\tau)Q(s_\tau | \pi) \end{aligned} \quad \text{A1.3}$$

Where $Q(o_\tau | s_\tau, \pi) = P(o_\tau | s_\tau)$ is the (predictive) likelihood of the (predictive) generative model

Remarks: Intuitively, the generative model of future states encodes beliefs that certain outcomes in the future are surprising (irrespective of the current state or policy), while future hidden states (given those outcomes) are surprising when they are not predicted. This generative model is defined in terms of a posterior (predictive) distribution over hidden states and a prior over outcomes. This contrasts with the usual construction of a generative model of past outcomes, in terms of a likelihood and prior over hidden states. Heuristically, this reflects the fact that current outcomes are caused by past transitions among hidden states but future outcomes can cause current state transitions (through policy selection).

Note that when $\tau = t$, the outcome is observed and the expected free energy reduces to the free energy of approximate posterior beliefs about hidden states:

$$\begin{aligned} G(s_t) &= -\ln P(o_t, s_t) \\ F_t &= E_{Q(s_t)}[G(s_t)] - H[Q(s_t)] \end{aligned} \quad \text{A1.4}$$

Optimizing this free energy corresponds to Bayes optimal state estimation; however, because this free energy functional has no concept of the future it cannot support purposeful behavior or active inference.

APPENDIX B

The variational updates are a self-consistent set of equalities that minimize variational free energy. Let $\tilde{x} = \tilde{s}_t, \tilde{u}, \gamma$ denote the hidden variables and $\hat{x} = \hat{s}_t, \hat{\pi}, \hat{\gamma}$ denote their sufficient statistics. Using the dot notation $A \cdot B = A^T B$, the variational free energy can be expressed in terms of its energy and entropy (with $\mathbf{B}(a_0)\hat{s}_0 = \mathbf{D}$):

$$\begin{aligned} F(\tilde{o}, \tilde{x}) &= -E_Q[\ln P(\tilde{o}, \tilde{x} | m)] - H[Q(\tilde{x} | \tilde{x})] \\ &= \hat{s}_t \cdot (\ln \hat{s}_t - \ln \mathbf{A} \cdot o_t - \ln(\mathbf{B}(a_{t-1})\hat{s}_{t-1})) \\ &\quad + \hat{\pi} \cdot (\ln \hat{\pi} - \hat{\gamma} \mathbf{Q}) + \beta \hat{\gamma} \\ &\quad + \alpha (\ln \alpha - \ln \hat{\gamma} - \ln \beta - 1) \end{aligned}$$

$$\begin{aligned} E_Q[\ln P(\tilde{o}, \tilde{x} | m)] &= E_Q[\ln P(o_t | s_t) + \ln P(s_t | s_{t-1}, a_{t-1})\hat{s}_{t-1} \\ &\quad + \ln P(\tilde{u} | \gamma) + \ln P(\gamma | \beta)] \\ &= \hat{s}_t \cdot (\ln \mathbf{A} \cdot o_t + \ln(\mathbf{B}(a_{t-1})\hat{s}_{t-1})) \\ &\quad + \hat{\gamma} \mathbf{Q} \cdot \hat{\pi} + (\alpha - 1)(\psi(\alpha) - \ln \hat{\beta}) - \beta \hat{\gamma} \\ &\quad + \alpha \ln \beta - \ln \Gamma(\alpha) \end{aligned}$$

$$\begin{aligned} H[Q(\tilde{x} | \hat{x})] &= \alpha - \ln \hat{\beta} + \ln \Gamma(\alpha) \\ &\quad + (1 - \alpha)\psi(\alpha) - \hat{\pi} \cdot \ln \hat{\pi} - \hat{s}_t \cdot \ln \hat{s}_t \end{aligned} \quad \text{A2.1}$$

Differentiating the variational free energy with respect to the sufficient statistics gives

$$\begin{aligned} \frac{\partial F}{\partial \hat{s}_t} &= \mathbf{1} + \ln \hat{s}_t - \ln \mathbf{A} \cdot o_t - \ln(\mathbf{B}(a_{t-1})\hat{s}_{t-1}) - \hat{\gamma} \cdot \nabla_{\hat{s}} \mathbf{Q} \cdot \hat{\pi} \\ \frac{\partial F}{\partial \hat{\pi}} &= \mathbf{1} + \ln \hat{\pi} - \hat{\gamma} \cdot \mathbf{Q} \\ \frac{\partial F}{\partial \hat{\gamma}} &= \beta - \mathbf{Q} \cdot \hat{\pi} - \hat{\beta} \end{aligned} \quad \text{A2.2}$$

Finally, we obtain the variational updates by solving for zero and rearranging to give:

$$\begin{aligned}\ln \hat{s}_t &= \ln \mathbf{A} \cdot \hat{o}_t + \ln(\mathbf{B}(a_{t-1})\hat{s}_{t-1}) - \mathbf{1} \\ \ln \hat{\pi} &= \hat{\gamma} \cdot \mathbf{Q} - 1 \\ \hat{\beta} &= \beta - \mathbf{Q} \cdot \hat{\pi}\end{aligned}\tag{A2.3}$$

For simplicity, we have ignored the derivative of value with respect to the hidden states (numerically, this simplification appears to make little difference in the Markov decision processes considered in this and

previous papers). Including this term leads to an additional term in the (Bayesian filter) updates of expected states corresponds to an optimism bias (Friston et al., 2014).

The variational updates for precision can be multiplied by $(1 - \lambda)$ and rearranged to give:

$$\hat{\beta} = \lambda \hat{\beta} + (1 - \lambda)(\beta - \mathbf{Q} \cdot \hat{\pi})\tag{A2.4}$$

This effectively slows the updates to provide a more time-resolved model of the implicit (e.g., dopamine) dynamics. In this paper, we used $\lambda = \frac{1}{4}$.